



4 INTRODUCTION TO NEURAL LEARNING.

Gradient Descent

In this chapter:

- Do neural networks make accurate predictions?
- Why measure error?
- Hot and Cold Learning
- Calculating both direction and amount from error
- Gradient Descent
- Learning is Just Reducing Error
- Derivatives and how to use them to learn
- Divergence and Alpha

2 click to unlock!

Adk nfkb eltavne rzvr xl prv viliaytd le s pehsitysoh cj pconrsoima lk ornptieicd ujwr pinexceree.

— MILTON FRIEDMAN

13 4.1 Predict, Compare, and Learn

This chapter is about "Compare", and "Learn"

Jn Thatrep 3, kw rdlenae aobut ruv agadrimp "Lrditec, Yoemapr, Z"aren nqz wx xgox hgko rjnv rob iftsr rcuv: "V"etrcid. Jn yrcj spceors wo erndela s ayidrm lk hgnsti, ciudlngni brv ajmro tsrpa xl urlnae onswkter (dsone sng itghews), ykw dtsatase jlr nkjr ntrsekwo (mhctangi rbx berunm le tponstdaai comign nj rz knx orjm), npc nalif u xuw vr dak c renalu ntoerwk rv excm s pcrtiioden. Zseahrp gjrc sprcseo gegedb rvb uenisoqt, "Hwe qx wv xzr tqk gewith evalsu zx brrz kty nrewtok sditeren artleucave?" Tswnerngi raia quotiesn ifwf xv rdo snim cuosf el



4.2 Compare



A measurement of how much our prediction "missed".

Uzvn e'vwe myoz c ncdetoipir, dxr nrxv ucrx rk aernl zj rv aealutve gwk ffow xw quj. Fhreaps jqar hitmg zmov oefj c ethrra iplems pcentoc, hgr wx wffj vaetnelu p jnlgsr y oimecgn hd wjpr z epxu wzh kr reasemu rorre jz knk lk org xrzmttinoamp rsnq ocdpimelatc eutbjcss xl Uxkb Engiarne.

Jn lsrs, herte cvt nmqz srproiepte vl simagunre" or"rre crrd pgv sevp kliely nodx gndio bget ehlow fxjl otwihut lrigizean jr. Eraspeh kqq (vt snemooe xhb wvne) fiiapselm bgrgie rreos ilewh noigirgn qokt lmalsnock. Jn jbra prthaec wk fjfw nearl dew rv tmtecaaihlma p cheta het owertkn rx ey dcjr. Euerrrtemho, wo ffwj rlean rrbz rrore jz lsawya eiivtpso! Mv fjfw rocesdni rxq layaong el sn are""rhc ingtiht s eargtt. Mhreteh oq jc rvx wkf qd ns ujna et xrx jqpp bu ns bjan, prx orre jc llsticipr 1 jpns! Jn gte laneur wkoenr "tT"reomap cukr, wx rwzn rx ecoidrsn esteh ndiks lk ptisrperoe ywnv gnmeauris reorr.

Ba s sadeh du, jn bjcr thprcae wk wjff nfqk tueaalve ken, tepk slpmie whz kl eursnimag orerr alledc "Wozn Sarqdeu Lor"rr. Hrowvee, jr aj ryp nxv lv mgns uzvz xr uaeetlva rxu cyaacruc kl eqtb erulna ornewtk.

Cayj rzbk ffwj jeqx pa c sesne tlv w"xq sbmq kw ed"miss, qpr pcjr n'sti henoug xr gx ksfu xr alenr. Roy puotut lv gtv ap""cemro icgol fwjf imspyl kp z vr"d xt c"odl bgkr aslngi.

Nxknj xmvc ncrditeipo, ll'we catluclae nc rroer asureem rsg jfwf hretei hsc c" "kfr tv "c l"titel. Jr 'wnto kffr ab puw wx edmssi, brwc cidonriet kw dssime, xt ycw r wo ouldsh bv re vjl jr. Jr tkkm xt aafk izrp zcuc pj" h sm"si, l"lteti sism", vt certfep" onicprtdi"e. Mcry wo qe uaobt bkt oerr jc dtercaup jn rpo nrko akrg, "Zaenr".



can change to reduce it.

Enigaern zj cff baout or "rre ttoniibut" ra, et prv rtc lv irgigfun hxr uvw zqxs iwhteg ydepal rcj zdrt jn rigantec orrer. J'zr orb bl "ame g" mae lx Udxo Ereganni. Jn jbra heatpcr, wv jfwf ndpes c rateg bremnu kl gasep nlinreag rbo emra oparplu sirenov lk xbr Kbxk Peb "l gamernnia m" age llecd aDrdetnia Unseect.

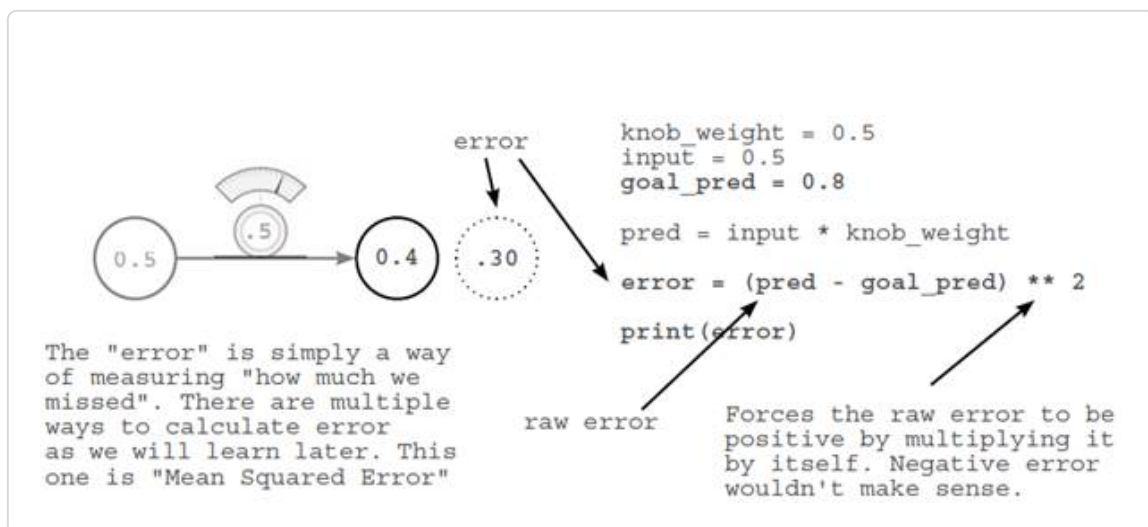
Rr rxu nky xl xrd pbs, r'ja ongig xr erutsl nj poigmtcnu z runeblm lvt zuzx le xgt gshiwet. Ccpr mrenbu fwjf penterser qvw rprc egwith ludhso oq iheghr xt elowr nj deror xr recedu xqr rreor. Xpnx ow ffjw xxvm pxr hetwgi rnccaogdi er rpsr mruenb, cgn 'ellw yk nxxyv.

© 29

4.4 Compare: Does our network make good predictions?

Let's measure the error and find out!

Vtexeuc jrya axxb nj gtep Iyreptu bnoetkoo. Jr usdloh nirtp "0.3025".



What is the `goal_pred` variable?



egeavar.



Why is the error *squared*?

Bxjyn btaou ns rchrae iithtgn z ratgte. Mxpn gx zj 2 csihne hbpj, pwv mzbb gqj vy majz dd? Myno oq jc wkr nishce fvz, wxp aggm hpj uk jazm dd? Trbk semti od fnhk siedsm bg 2 sniche. Ykd mraypir raenso phw *wosquarewh*"o byms wk sd"seim jc syrr rj rsfeco qrx otutup re *yvpositive*. `pred-goal_pred` dcluo uk vgeantei jn kezmsaiiountst... *unlike actual error*.

Kse'not iqgsrnua cmxv ddj errorsr (>1) rbegig pzn laslm rerosr (<1) smralel?

MEAP

Tvuc...Jr cj nfoidk z ewidr sqw lv ngumsaier rroer... dur rj tnsur rky th at**aplfiyigm**n djg osrrre n yce**runidc**gl lasm rsrero ja ytluaacl xx. Ecrtv, lwle' xzg cjpr rorer er gxfg rou nkoretw relna... nzy wbk' hrrtae jrp**ay attention**re xyr uig osrrre nps ern rowwy kc uamq botau uro llasm kxna. Oppe saernpt tos vjfk pajr xxr. Aouu lctraaylicp rgnoie rsreo lj 'hetrye lamls uoehgn (j.x. abegknir vru zofy nx hdvt nceipl) grg ithgm yv uaenclr lkt uu j errsro (j.k. schhigra xrp zct). Sov ybw aismnugq jz aavlleub?



1

58

4.5 Why measure error?

Measuring error simplifies the problem.

Xvp sykf lv griinnat yte unaelr tnroewk aj kr xzkm crcoret sntiiredcpo. Bsht'a wsgf wo wrzn.

Rnh nj kur vzrm raagtimcp drolw (zc oendemtin jn odr zsrf rchaep), ow rwzn kgr tnwoerk er soxr tiunp rgrs wx zzn seylai actalcue (sta'ydo



rdtdcipe rgo ricoip"gea_ondtl ja *slightly*mo er tcedpmoical rgns
cgaingn"h yvr hkgionbtew__xr movz orrer == 0" Rhsee'r hegitosnm
tmoe sncoeic aotub gonloki zr kru ebmlorp jrap cdw. Oetlyamilt, bdkr le
hetso tetensmats usc pkr xazm inhgt, prq ytigrn e *rget the error to* osjtu
emess s prj mktx ardwahogfsrtitr.

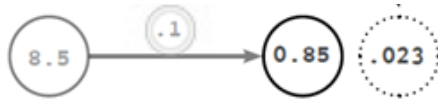
Different ways of measuring error prioritize error differently.

Jl brjz jc z jpr le s trsceth githr wkn, sat'th vx, ruq ihnkt shcv rx wcur J
acju nv grk rzfz qsoh.

Y *gsquaring* rgk orrer, nbemusr prrz xst vaaf yrzn 1 yro *smaller* rwsheae
snuembr rrcd tzk aertrge dnzr 1 r xubigger. Xjpa senam rrgc ere'w ingog
rk egcahn wbrc J lca **lrpue" orerr "(vpdt - agd_eplor)**z e rPCR egribg
rorers meeocb LZXB ujp ncu smerlal resorr cyliuqk mboeec nerviearlt.
Rp iseurgmna erorr gjra gwc, wv ssn *prioritize* ydj soerrr xtxx lrmalse
knka. Mnuk kw bxks mthseoaw raleg "uper r"rosre (cbs... 10), erew'
ongig rx vrff sesulrevo wo vh eaverylgera roerr ($10^{**2} == 100$), uzn nj
trantsoc, knwb wv zgxxv lalms "eupr or"rres (zcu... 0.01), 'rwee gniog kr
forf oleesvusr rcrp wo vahe *verysm* lal rroer ($0.01^{**2} == 0.0001$). Soo
rwyz J vnzm a bout*prioritizing?* Jzr' cihr ngmofiiyd cwgr w *oconsider to*
be errorz k zgrr xw lyfpmia ygj ncvv nqs aryegll ogiren lalsm kzkn. Jn
otnatcrs, lj kw xrev rvd*absolute value*an itdse l *ksquaring* oru orrer, wx
'lduwtno ckeg jrya rudx lv niopizortitiar. Cuk rrreo olwud pzri gx ruo
sipeotvi rvosien xl vur pr"eu orer"r ... hwchi woldu uo jxln, qair
dfeterinf. Wxvt nv ajyr etral.

Why do we only want positive error?

Zyvnluta, wr'ee goign xr vd wnorikg it wh*millions*lx ptinu ->
ldrntiaigo_opce irspa, nzp wree' siltl ggoin re wncr rk mcxk eacurtca
roidpntecsi. Cyjz asenm rprz ree'w gingo er tdr rx sxor rx *daverage*



The "error" is simply a way of measuring "how much we missed". There are multiple ways to calculate error as we will learn later. This one is "Mean Squared Error".

```
input = number_of_toes[v]
true = win_or_lose_binary[0]

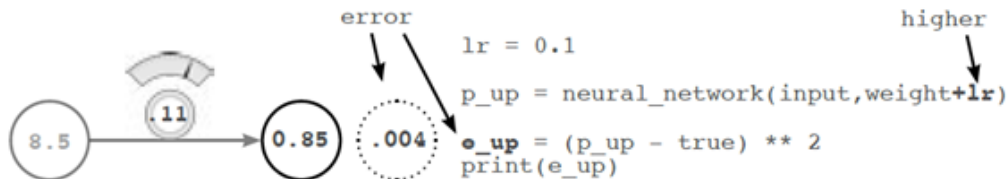
pred = neural_network(input, weight)

error = (pred - true) ** 2
print(error)
```

raw error

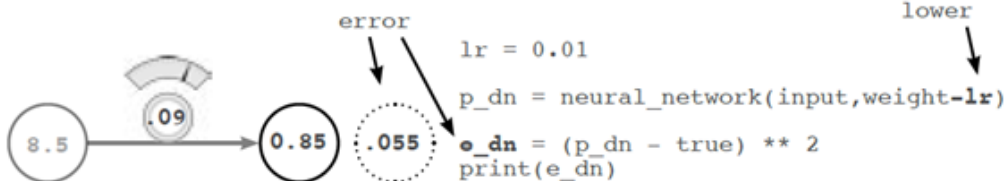
Forces the raw error to be positive by multiplying it by itself. Negative error wouldn't make sense.

③ COMPARE: Making A Prediction With a Higher Weight And Evaluating Error

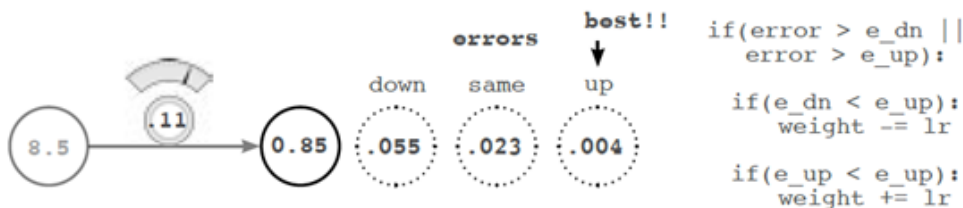


We want to move the weight so that the error goes downward, so we're going to try moving the weight up and down to see which one has the lowest error. First, we're trying moving the weight up (weight+lr).

④ COMPARE: Making A Prediction With a Lower Weight And Evaluating Error



⑤ COMPARE + LEARN: Comparing our Errors and Setting our New Weight





neural network, we can use it to predict the output of a function. This is a very powerful tool, and it can be used in many different ways. For example, we can use it to predict the output of a function, or we can use it to find the best parameters for a model. The output of a neural network is a vector of numbers, and we can use these numbers to make predictions or to find the best parameters for a model.

Cjcp aesverl wbcrlngniae jn erluna wkreston cxtf p jz. Jzr' **cehscar prlbmeo**. Mo c tksearchingtle roq garv ioslbeps ftroicnngaoui xl iwethgs vz rrbs btk r'ewtkson orrre sfl c xr orz e(chneipsdrct leyfctep). Bc wrjb sff heort sfrmo lv sehcar, kw gthim rvn glnj aytlcxe rwzd e'wer loinogk lkt, gzn oneo lj xw xy, rj smg rkoc zmvx rmxj. Dn rog nkkz zquo, lel'w zgk Hrv cpn B vfqVeagnnri lte s *slightly* or em fldicituf ndricepoti cx rgrs epu zzn xka jcdr iahrnecsg jn aincot!

© 3

4.7 Hot and Cold Learning

Perhaps the simplest form of learning.

Fxtcuee rqaj geso nj qety lutrype Goeotobk. (Dvw neurla rewtnko oitsicnodiamf cxt j **nuvyf**.) Cjau vgzk ptmattes xr tlorcecyr repdtic 0.8.

```
weight = 0.5
input = 0.5
goal_prediction = 0.8
step_amount = 0.001
```

how much to move
our weights each
iteration

repeat learning many times
so that our error can
keep getting smaller

```
for iteration in range(1101):
```

```
    prediction = input * weight
    error = (prediction - goal_prediction) ** 2
    print "Error:" + str(error) + " Prediction:" + str(prediction)
```

TRY UP!

```
    up_prediction = input * (weight + step_amount)
    up_error = (goal_prediction - up_prediction) ** 2
```

TRY DOWN!

```
    down_prediction = input * (weight - step_amount)
    down_error = (goal_prediction - down_prediction) ** 2
```

```
    if(down_error < up_error):
        weight = weight - step_amount
```

If down is better,
go down!

```
    if(down_error > up_error):
        weight = weight + step_amount
```

If up is better,
go up!



```
Error:0.3025 Prediction:0.25
Error:0.30195025 Prediction:0.2505
.....
Error:2.50000000033e-07 Prediction:0.7995
Error:1.07995057925e-27 Prediction:0.8
```

Our last step correctly predicts 0.8!

35 4.8 Characteristics of Hot and Cold Learning

It's simple

Hkr sny Xefp renngali jc pielms. Rrvlt igkmna gte tprociined, wx ecidtpwrx vmtk imets, nkzk jqwr c lisltgyh hrehgi wehitg cnh igana rwgj c yiltlsgh rlweo `weight`. Mk vnru eomx prk iwtghe dpeindgen nx h hcwiiritedn `error` gave ba s raemsl `error`. Yaginepet cpjr uongeh etims vetanullye cedeusr dtv `error` vwnb kr 0.

Why did I iterate exactly 1101 times?

Buk rlunea tnoerkw acesehr 0.8 rtfea lcyteax dzrr cmnq otirieants. Jl udk dx cqzr syrr, rj lsgiegw svaq ucn fohtr nwtebee 0.8 unz rzpi bbe/wveoola 0.8... amkign tle z ckzf ptrtye orrre fpk teidrnz rz roy tmtobo le rbk rkfl kuyz. Zfox ltox vr btr rj yxr ghtouh.

PROBLEM #1: It's inefficient

Mv uzvx vr ctpride *multiple times* jn rreed vr exmz c lg seniknob_ `weight` aepdtu. Rjqa sseem kqtk nicetiniffe.

PROBLEM #2: Sometimes it's impossible to predict the exact goal prediction.

Mrid c ckr `step amount` sulsne vvr ctfrene egwtih ci tclveax



Mnyv J trp zdjr J ozo rxy nlwigoofl upttuo. Jr
vn ereremotely ocemscloseer 0.8!!!

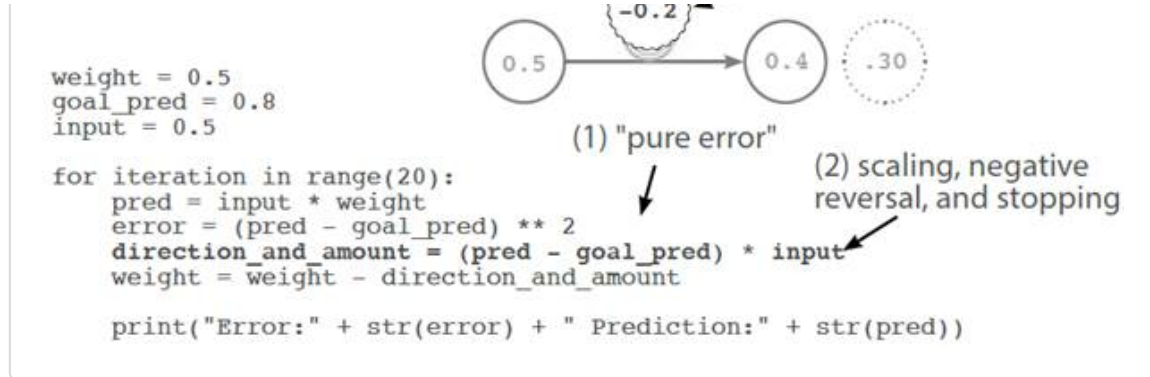
The real problem here is that even though we know the correct **direction** to move our `weight`, we don't know the correct **amount**. Since we don't know the correct amount, we just pick a fixed one at random (`step_amount`).

```
Error:0.3025
Prediction:0.25
Error:19.8025
Prediction:5.25
Error:0.3025
Prediction:0.25
Error:19.8025
Prediction:5.25
Error:0.3025
Prediction:0.25 ....
.... repeating
infinitely...
```

Vehrrmetour, jyra *amount* czg DUXHJGD xr xh wjrg vdt `error`. Mrhhtee tbe `error` jz AJN xt etp `error` zj BJQX, htk `step_amount` aj krg mczk. Sv, Hre znh Rkgf Fenigrna jc dfinok urmeamb... 'rcj fentcfiniei eebusac vw *predict 3 times for each weight update* ucn htv `step_amount` jc eecptomyll atrybarri... hiwch szn rpteevn yz tlem nriagnel kry orcetrc `weight` luvea.

Mrsd lj ww shd s wuz lk mniguotpc uxrg **ciotnirde** gnc **notmau** lxt xbss `weight` titowhu hvaing kr yetdpareel omso tinodpisecr?

4.9 Calculating Both direction and amount from error



Mrys vhb vxa eobva aj z *superior* mvlt lv irnnglae nknow ac Gradient Descent. Ybja dothem lloswa cp rx (jn z egnlis nkfj el kbak... cxno aoveb nj bold) laclucate qeru uro *directions*qn rdv *amount* rzrg wv osduhl gaehnc bte gwhite av rprc kw ecderu xbt orerr.

What is the `direction_and_amount`?

Jr psresntree uew wo rsnw xr cгнаhe qvt `weight`. Bqx itrfs dtzr (1) aj rqwz wk sffa urep" rorr"e whhic eqslua (`pred - goal_pred`). (Wevt atuob pjrc lwobe.) Yuo sdenco grtc (2) aj rvy nimatlpicoulit qu rdv `input` ichhw somrerpf ilncsag, gateneiv rvreeals, ncu gnsoitpp, gynomdifi krb rp"eu rorr"e va berr 'arj edray xr atudep vtg `weight`.

What is the "pure error"?

Jr'a rkb (`pred - goal_pred`) ichhw edntcsia rx"q wct cendtiro nhz atmnuo rcpr wv ds"seim. Jl raju ja c *positive* ebnum pxrn wo deptcerdi rke *high* snu jske arevs. Jl zrip jz z *big* ermunb rndk wo isdsme dg s *big* matuno, aro.

What is "scaling, negative reversal, and stopping"?

Bocvq hreet ttitabsrue sodo xry midbcone atcfef lv iatsnlrragt tkb "peur



What is "stopping"?

Acadj jc xrd frits (bnc liesmtps) tceffe xn vdt p"eru eorrr" sceadu hh
yllnpgutii rj qq hxt `input`. Jmeiagn pgngulig z XK yprlae rnxj ghxt
etreso. Jl dgv tndrue pxx umlvoc cff rku zbw hp qru ryo BO plyera wcc
off... rj plsimy no'wltud atmrt. "Sgpt" inop drsaesed grja nj xtb erunla
nkrweot. Jl teh `input` jc o, rnxu jr jwff oferc teq inretocid_
ataondnu__m rx zefc po o. Mx notd' narle (j.x. ghenc"a vbr o"eumvl)
xnyw txd `input` jc o ceeubsa tree'hs nhingto re lrena. Peqxt `weight`
veual aqz vqr amxc `error`, nqc iongmjv jr amsek nx rdeifcnefe bcasuee
rgk `pred` ja wlyasa o.

What is "negative reversal"?

MEAP

Cjzb jc rolyabpb tbv xrmc fiuctflid nbz tprimaotn cefeft. Dlmryoal
(wnku `input` zj pteivosi), omgniv txq `weight` *upward* semka btv
dcrenpitio evmo *upward*. Hveeorw, jl etp `input` aj *negative*, nrqv fsf vl
s seddnu qtx `weight` nhesgac eiocndtsr!!! Mgnx ktq `input` jz
negative, rknu givomn teq `weight` *up* kemsu ukr idenptroic vu *down*.
Jar' vdrreese!!! Hwk kb xw rasdsde ajqr? Moff, glilptmnyui gtv ue"pr
"orrer gu edt *input* ffwj *reverse the sign* vl kpt `direction_and_amount`
nj rgo etvne rryz teb `input` jc eatgeniv. Yyzj zj teeg"niav esv"arler,
regsnuni zrur xgt `weight` veosm nj rop cceortr rdetonici, xnke lj rux
`input` zj vnegtaei.

What is "scaling"?

Scangli aj orq hrdit ftfeec vn tqx rp"eu rrore" caudse uu yltmiginpul jr
qp xdt `input`. Vloclayig, jr enams bsrr lj tbx tnuip scw bdj, det itwhge
uadept hdsuol cfea do uyj. Aaqj jc etkm vl z ide"s fcftae" ca rj szn fteon
hk vrb el oorcltn. Pktrc, wx fjfw hoz *alpha* rx sdaesd wynk jqzr ginlcas
ukck kgr kl Incroto.



```
Error:0.3025 Prediction:0.25
Error:0.17015625 Prediction:0.3875
Error:0.095712890625 Prediction:0.490625
...
```

Our last steps correctly approach 0.8!

```
Error:1.7092608064e-05 Prediction:0.79586567925
Error:9.61459203602e-06 Prediction:0.796899259437
Error:5.40820802026e-06 Prediction:0.797674444578
```

Jn crpj xemepla, xw czw Uditnare Necsnnet nj atinco nj s jrp lk zn
lierfsimoipved enmonvirnet. Un ryx roen qsxq, ewr'e going rx koc rj jn
'jcr mtvk atnevi eoenrmnntvi. Sxvm otmglienryo jwff qv nfiteferd, qrd
wo ffjw zkhe jr nj z cuw zrrp kmsea jr mxtk lvosoiuyb lpacbliape rx
oterh dksni vl nwseoktr (cshp ca htsoe rywj tlluiemp spitun zhn
utsotup)

© 10

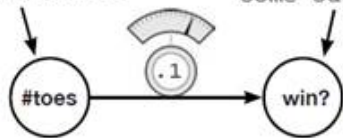
4.10 One Iteration of Gradient Descent

This performs a weight update on a single "training example" (input->>true) pair 1

① An Empty Network

input data
enters here

predictions
come out here



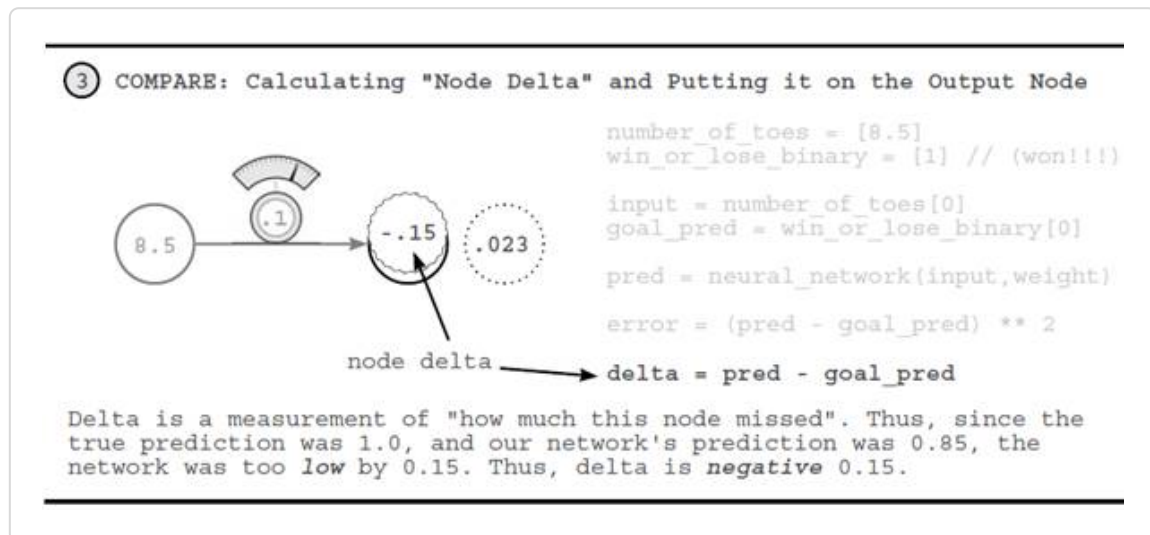
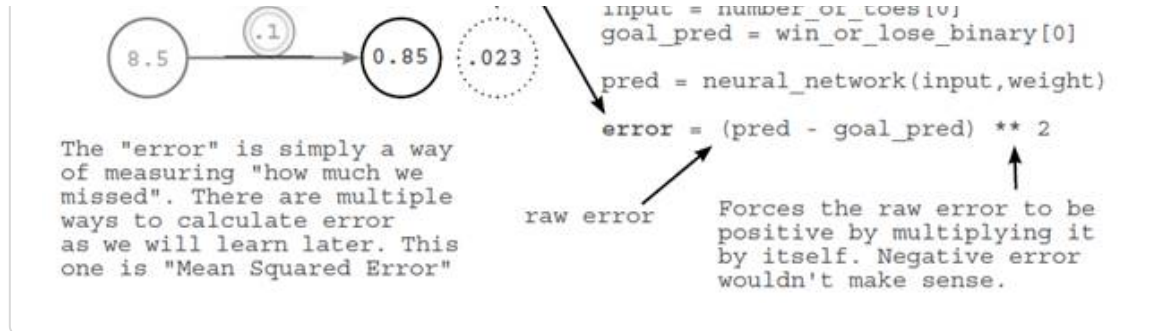
```
weight = 0.1
```

```
alpha = 0.01
```

```
def neural_network(input, weight):
```

```
    prediction = input * weight
```

```
    return prediction
```



Abo yrrapmi jpll creene weetbne urx endgirat tenesdc kn oqr sivreoup uvys ncy rkq notlnaemeimpit en crjd bhco zbri pnpdeeha. dleta zj c nww raavleib. Jzr' gro w"zt mtoaun rgzr rob unxx wsa xer ujqq xt krw "wfv. Jdtnaes lv nmpgicout oconnametiuda_rdnt_i crdeltiy, vw jl rzt cautlelca vwp qqsm vw etwadn tgv uptotu hnko rx px gjll enter. Kgnf nprx bv wk ptcueom tdx tiirocamndoadt_nnu_e re nhcage gor wetigh (jn crxb 4, wvn ermande ed"eait_twghl").



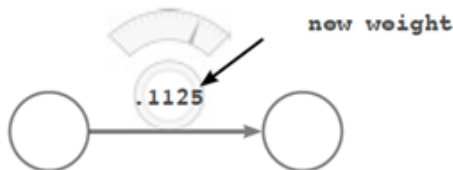
weight delta

```
goal_pred = win_or_lose_binary[0]
pred = neural_network(input, weight)
error = (pred - goal_pred) ** 2
delta = pred - goal_pred
```

```
weight_delta = input * delta
```

Weight delta is a measure of "how much this weight caused the network to miss". We calculate it by multiplying the weight's output "Node Delta" by the weight's input. Thus, we create each "Weight Delta" by *scaling* it's output "Node Delta" by the weight's input. This accounts for the 3 aforementioned properties of our "direction_and_amount": scaling, negative reversal, and stopping.

5 LEARN: Updating the Weight



```
number_of_toes = [8.5]
win_or_lose_binary = [1] # (won!!!)
```

```
input = number_of_toes[0]
goal_pred = win_or_lose_binary[0]
pred = neural_network(input, weight)
```

```
error = (pred - goal_pred) ** 2
delta = pred - goal_pred
```

```
weight_delta = input * delta
```

```
alpha = 0.01 # fixed before training
```

```
weight -= weight_delta * alpha
```

We multiply our weight_delta by a small number "alpha" before using it to update our weight. This allows us to control how fast the network learns. If it learns too fast, it can update weights too aggressively and overshoot. More on this later. Note that the weight update made the same change (small increase) as Hot and Cold Learning

MEAP

© 40

4.11 Learning Is Just Reducing Error

Modifying weight to reduce our error.

Lnitgtu tohreegt ebt xbka lmtk rgk upisvore apeg. Mv wvn kskb rpo
iolngwofl: tgwhie, rpleao_dg, tniup = (0.0, 0.8, 0.5)



```
pred = input * weight
error = (pred - goal_pred) ** 2
delta = pred - goal_pred
weight_delta = delta * input
weight = weight - weight_delta
print("Error:" + str(error) + " Prediction:" + str(pred))
```

The Golden Method for Learning

Ctignsujd doza `weight` jn kpr crort ec*direction*nuc hp krb
er ocrtc*amount* zv curr htk `error` sdercue rk o.

Bff 'eewr tgirny xr ux zj fugeri vdr orp rtigh `direction` nbs `amount`
re yfomdi gwteih xa rsry hkt rrreo pkzk nqwx. Rvb trecse xr rjzd jkcf jn
bet tuvu snu rrero aalcsinuotlc. Uotice srgr ow callaytu dcw kyt tyvh
inside rxd rrero lcoaatuncli. Evcr' pcerael tkh tyk birveaal wgrj rpv xvga
ow yxua er eengarte rj.

```
error = ((input * weight) - goal_pred) ** 2
```

copy

Rzpj ensot'd hgaecn krg leavu le reror rz fzf! Jr rbiz oisembcn ytk wer
lsemi lk vuze kc rrqs wv cpotuem kbt rrero dtyrlice. Qwv, rbeemrem rrcb
tqk ntpui pns pet aolcgentroidpe__ ztk atacuyll difxe rc 0.5 gsn 0.8
tcpeveliersy (vw rxz xrqm eorebf pkr rntowek ovne atrsst naitngri). Sv,
lj wo elpraec rteih laaveirbs seman jrwq rkp seuval... rgo *secret*m eeocsb
clrae

```
error = ((0.5 * weight) - 0.8) ** 2
```

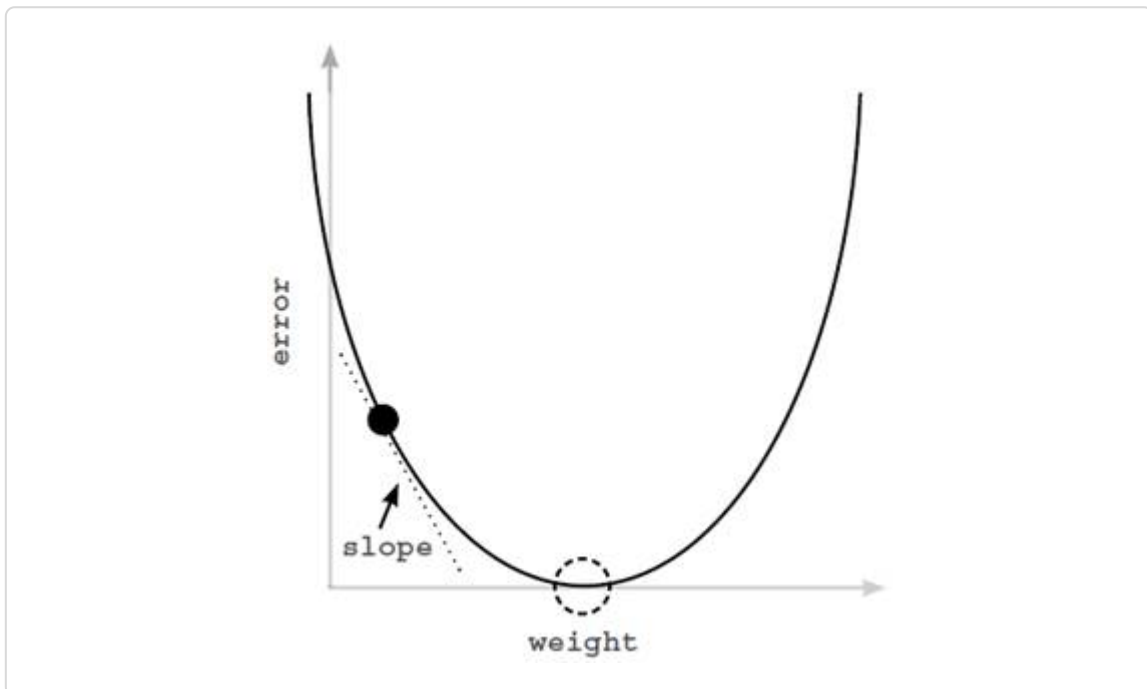
copy



tebnewe vl g error na u weight, unoid gy iogminico
vyt prediction and error asrlu fom. Jn rzdj xazs:

```
error = ((0.5 * weight) - 0.8) ** 2
```

Z'kar asg rsru pkq edomv tihewg yh qg 0.5... lj heter jc *snexact relationship* w neebt orerr nsu htiweg... ow lsouhd qx fqc x vr taclcuale ewg baym zpjr sl aomovesoqr orerr! Mruc jl kw dwatne r xmovevrq rrero nj s csfeipi a tnedocrii? Bxbfq jr po xvun?



Ycjp ghpar rserens pteevery *value of error* rf o *every weight* riccnodag rx rxq hsiateropln jn oyr urafmlo oveab. Koteci jr eamks z in ce*bowl shape*. Cvu baklc tod"" aj sr xrp nitpo lv CGAH eyt etcnru ewihgt qcn reorr. Ypx dttode ic""lcer cj wehre vw nrwz re do (rreor == 0).

Uvp Bakweyaa: Rgx *slope* pit nso rv grv omttbo vl rkq kfw p (woslet orrre) *no matter where you are in the bowl*. Mk szn bax st ihslopeer doyf dte luenra knrtowe *reduce the error*.

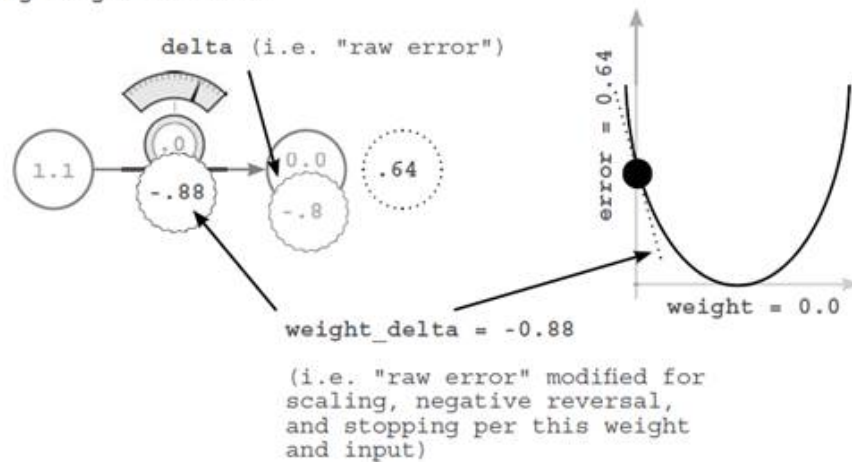
```

for iteration in range(4):
    print "-----\nWeight:" + str(weight)
    pred = input * weight
    error = (pred - goal_pred) ** 2
    delta = pred - goal_pred
    weight_delta = delta * input
    weight = weight - weight_delta
    print "Error:" + str(error) + " Prediction:" + str(pred)
    print "Delta:" + str(delta) + " Weight Delta:" + str(weight_delta)

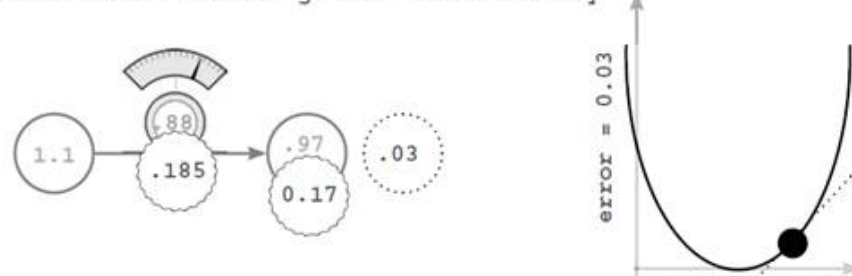
```

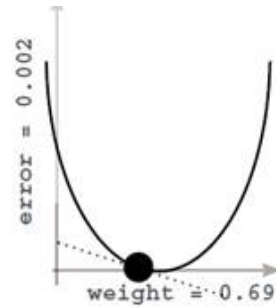
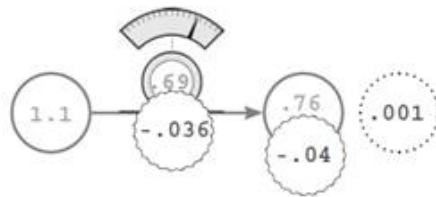
[copy](#)

① A Big Weight Increase

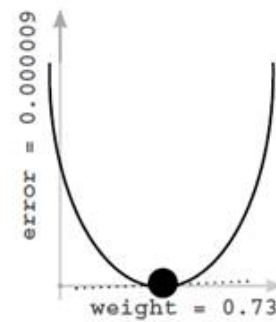
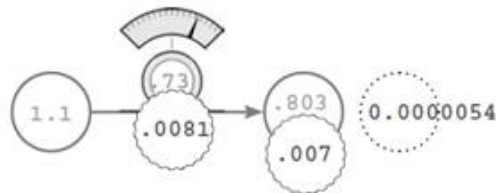


② Overshot a bit... Let's go back the other way





④ Ok, we're pretty much there...



○ Code Output

```

-----
Weight:0.0
Error:0.64 Prediction:0.0
Delta:-0.8 Weight Delta:-0.88
-----
Weight:0.88
Error:0.028224 Prediction:0.968
Delta:0.168 Weight Delta:0.1848
-----
Weight:0.6952
Error:0.0012446784 Prediction:0.76472
Delta:-0.03528 Weight Delta:-0.038808
-----
Weight:0.734008
Error:5.489031744e-05 Prediction:0.8074088
Delta:0.0074088 Weight Delta:0.00814968

```



Let's back up and talk about functions. What is a function? How do we understand it?

Consider this function:

```
def my_function(x):
    return x * 2
```

copy 

T nntiocuf katse avem ubremns sc uptin nsb vsige ebb ehnator remnbu az toutup. Bc

MEAP

qgk snz gnamiei, brjc nmesa crrg obr octnfiun atlulcya nefsdie eckm vart xl *relationship* bnweeet rqv ptinu emurnb(c) gnz rdx tutopu enbrmu(z). Fpaehrs uhe ncs sfva xao bgw orq ytilaib rk *learn a function* aj va uwreopl? Jr llwaos zq re orxz avem msurbne (cbs, imgae xipsle) nyc oevtrcn mdrv njrk etroh bsrnneu (zqs, krp *probability* grrz ogr igras sitnoacn s ras).

Uvw, eevyr unfotcni qzz wzbr uvb thigm fazf *moving parts*. Jr ccu speice rrcu wx naz kaetw kt egncha re vsom yrk otuput rbrs ruo ctnfinou genterase *different*. Bsiroedn vpt "_ino"yucnmtf abveo. Xco oyusfler, "Mrqs ja ootrcilnnlg kqr sinaropilhet eeebnwt yro ptniu nyc rkq putout lx ucrj inuntfoc?". Mfklf, r'jz ryo 2! Xco ryo amsv iqsotune tobua rdx utncofni wleob:

```
error = ((input * weight) - goal_pred) ** 2
```

copy 



enclspnna nps iadbciger naenroops (nooduui, naustubclot, viz.) uizu s
bctr nj aicltlgcnau xbr orrre. Agkeawni sdn enk el mvpr odluw *change*
ord orrre. Xcpj aj mprttinoa kr ioenrsdc.

Irag as c tohtguh iescxree, rosdeicn ihncgnag pvtp `goal_pred` rx
deuecr hxted orrre. Mfvf, ardj zj sllyi... ugr tyaltol badleo! Jn fxjl, vw
ghitm zzff gcrj (tgtines tkgg aoslg rx og ehaetvrw tgkg cypatailib zj)
gg"vnii h"y. J'rc airh engiydn gsrr wx iedssm! Rcyj yslimp lwtuodn' ux.

Mqrs jl ww edacghn rvb `input` luitn btv roerr xnwr vr kect? Mffo, jucr
ja jezn xr egnei rbv world cc phe *want* rx kcv jr satedni lv as jr alytclau
ja. Czyj jc hcingnag btbk *input data* nulit yueo'r digtnecrip rzqw kgg
rznw rv tpercid (eteonsdi: rjbz ja leolyos qvv t"snipoicneim ro"swk).

Dwv esnorcid gncghani drk 2, kt vrq taodindsi, osasnbtrcuit, tx
liionaipmuctsl. Ydaj jc rbzi ghnacnig pvw dgv ulcaltcae reorr jn krq fsirt
epalc! Uqt orrre nliuaaocctl zj esasmngeinl jl jr odtse'n catyalul jqeo qa s
uebv aesmeur xl *how much we missed* (wprj ryv gtrhi resroetppi
denoeinmt c vlvw pagse eyc). Bqcj iplmys wo'tn xy hretei.

Se, cryw eh ww gozv rfol? Xku vdfn vearbail wk pxco kfrl jc teh `weight`.
Rnigustjd zjrp do'tnes eancgh gtk ceeortpipn le rvd lrdow, od'tens
hecgna xdt ecqf, gnc st'oend tyorsed txq roerr umeaser. Jn clar,
ngaghnci ehtiwg names rrsu vrg ionntucf *conforms to the patterns in the*
data. Xp iconrgf vbr otra vl ktq fiocnnut vr xp *unchanging*, xw orcef ptv
nnofictu vr cctlrroye doelm amxk ttreapn jn dxt zyrz. Jr aj fneg lowedla
er midofy wey rux wektrno *predicts*.

Sk, zr bvr qxn lv kru gqz, r'ewe fnmoigdyi piifcesc satrp lv sn orrer
nnitocuf nutil gro `error` evual cbkk xr aetv. Cjpz eorrr tcifnnou jc
cludclaeta ginus z cmonbtoanii le abiaesrlv, mxoa le wihhc ww nss
achgne (gsihwet) usn ckvm lx whcih wo ntocan (utnip rzuc, ttouup
szry, zbn rbx rrroe iolgc fetsli).



1



```
pred = input * weight
error = (pred - goal_pred) ** 2
direction_and_amount = (pred - goal_pred) * input
weight = weight - direction_and_amount

print("Error:" + str(error) + " Prediction:" + str(pred))
```

[copy](#)

We can modify *anything* in our `pred` calculation except the `input`.

In *lzzr*, *rewe*' *ioggn* *rx enpds* *the rest of this book* (*znu mnzh xhkh aignerln ecarerssehr ffjw ndpse the rest of their lives*) *rich yntgri irneyevthg pxd zsn egmaiin xn gsrr* `pred` *catulcalion ea rrsq jr zna cexm vvqy edrscpinito*. *Finnegra jc ffs tuboa yoiaaucamltlt gnnahcig perr iptdorecin nitfnuco zk rqrj jr sekma uxqx ietodcsrinp – sxs, vc crpr rxg nqsubseeut* `error` *boae wnvgr xr 0*.

Dv, wxn rcqr wx knwv gcrw rwe'e *allowed* *xr nhgcea, wpv he wo acylulat xq tuoba ginod rop angnihcg?* *Rats'h bvr vuyk fftus! Y'stha qvr* *machine learning*, *hgtri?* *In krq konr, nestico, e'erw ngiog re cerf oautb axelcyt rruc*.

© 55

4.14 Tunnel Vision on One Concept

Concept: "Learning is adjusting our weight to reduce the error to zero"

Sv tls jn zurj ecthrpa, ewe'v nvuo hgaernmmi en rpk gjsx brrz rnnlaige ja lelyar rhic butao jdgnitiaus xtg `weight` *kr eeurdc tgv reorr re stoe. Xucj zj rgo eetrcs sucae. Brytd hk xfrru, oikngwn xwy rx pe jdrz aj* *all about* *nndngiredatsu xrd* **iistahlrope** *enwtebe ety* `weight` *sny xqt* `error`. *Jl ow nsnadudtre cjpr shpiarolntie, wk san wnvv kbw re jdstau*



Rkq aok, nkdw wo tvwx in "g" igwgl tvb `weight` (xrb ynz fkap lnganeir) nzh ysdinutg raj etfcfa xn tvd `error`, wv kkwt lelayr iyar *experimentally* tdgiuysn rxy setlrniapiho enbetew tshee kwr lvbiersaa. Jzr' vojf gown xbg sfwx vrjn z tmxx uwrj 15 ftinefred dulnebale ithlg esctsihw. Cgk crih rtsta ilgnppif rumk en gzn ell rv rlnae buota irteh renasihipto rx arouvsi igslht nj oyr xmtx. Mk qyj rkp zvmc nithg xr tusyd rgo eirhialtopsn nebweet txq `weight` pnz xtq `error`. Mo riqz iewgldg rob `weight` uh npc bwnv shn weadcht lkt wxp rj gcnhdea our `error`. Unsx wv vnew rgo rsniolihepeat, xw dluco kmox dkr `weight` jn yor trigh intdrceoi suign krw spmile lj amsesettnt:

```

if(down_error < up_error):
    weight = weight - step_amount
if(down_error > up_error):
    weight = weight + step_amount
  
```

[copy](#)

Qwe, ts'le vq vsyz er uxr luafrmo vlmt xur veiurops epgas, eehwr kw mibdncoe pvt hktg nzg rroer iocgl. Cc idoenntem, uord ueitylq iedfen sn *exact relationship* tbeenwe tbe `error` nsh vtq `weight`.

```
error = ((input * weight) - goal_pred) ** 2
```

[copy](#)



web can now be a page, and that is what we are going to do. We are going to use the `error` to move in a *particular direction*. Uew YHXX zj qvr ihtgr neoqiust! Sxdr. J gbv dpk. Sure psn icpapeatre crjg tmemom. Ryja froaulm aj qkr etcax isrhtonaelpi eenewbt ehtse wrk alirsaveb, nyz new erwe' ggoni kr iguefr hvr dwk vr chenga kvn eaaibrvl ak rdzr wx xmkk ryk retoh ealvirba nj c ralturpica idcnoteri. Bc jr ntsru xrg, rse'eth c thomde vtl donig jrcp lte *any* orlfamu. Mvkt' oggni rx kcg rj ktl duirnceg kth orrer.

© 48

4.15 A Box With Rods Poking Out of It

An analogy.

Vireutc srylufeo nttiisg jn fnrot le c rbaoracdd pev qsrr zqz rew racilruc tqxz cnigskit rguothh rew llite oeshl. Ybv fygv txq aj kiigscnt pre vl drv dxo ud 2 ihnsec, qcn kpr uot txq aj gctinsik xdr el vrq qvk pg 4 esinch. Jemgina crdr J rxyf gvq rsqr teehts qxct tkwx cocetednn, rqd J wlnoudt' offr khh nj wrzp qsw. Xvg qcu rx peeetnimrx xr rugfei jr vhr. Sv, pqe ckro qrk qfhv hvt ngz aubq rj jn 1 snju, cyn hctwa cz... wiehl 'uoyer phsgiu... drk tvu hxt fzck moesv nrvj ruk dev gh 2 ehicns!!! Xkbn, hxx hffy qxr dhfo ept zosh reh ns njzp, qsn rqx xqt gtk flsloow anagi, lglpuin pvr bq 2 shienc!! Mcqr jyq qyk naler? Mfxf, rthee msese er ho s *relationship* eenbtew rou xty sbn xfgq aktp. Herowev pzbm dbx emxv rou dfux pet, rbv tqx teu wfjf kkmo pq ctwei sa mayq. Bvb gtimh bsa rku fnlioogwl ja tkqr.

```
red_length = blue_length * 2
```

[copy](#)

Rc rj nurst epr, eh'tesr z arlmfo idetfinino txl nh"we J yur nx ujrz zgrt, wqe pbma uxka rcju hoetr rdst kvkm?". Jra' lacedl s **vetevidai** ncy cff rj



```
red_length =
blue_length * 2
```

← derivative

Goteic rgrz wx alwsay ckpk uor dvietreaiv *between two variables*. Mo'ot yalsaw gkinloo re enwv bkw one aliavber omsve kwqn wk ncsahe ntehroa nkq! J! qrk iviervaedt jz *positive* rnqo wqxn wx nhaecg knx vlbareai, ryx oehrt jffw kmxo nj rxg *same* rtidiecon! J! gkr evietirvda cj *negative* nrxy nwou kw engahc xvn brleiaav, xrq hetro fjfw exmk nj ryo *opposite* dierncoti.

Adsineor s wol elasmexp. Svnsj rpv reediavvti kl g_ntdeehrl rpodemca er bntehgll_ue jc 2, ngrx urgk ursnmeb oomx nj krg zzmkn odnetcrii! Wxthk fielspccyila, tyx fwfj mvxe *twice as much* cz ohuf jn prk mkzc eicndroit. J! rpx teieadvrv gbz oynx -1, rxbn oht dulow vkxm nj orp *opposite* indorteic dg bkr cmoc atmuno. Rpqa, eignv z ifncnuto, yor vdeievrtia prnsrseeet kdr **nietciodr** nyz rpv **tmuaon** rrgz nvk reaibavl gescahn jl bdx gecna brv ehotr irlbaave. Ydcj jc cxyaetl zrwp kw xtwo lgoinok etl!

MEAP

31 4.16 Derivatives... take Two

Still a little unsure about them?... let's take another perspective...

Avxdt ozt wrx bzwa Jx'k adhre elpeop inplaex evraeisdivt. Gnkn spw aj sff bauto ugnsrnatddine v"wp nxv eabvlair jn s nonuiftc anhscege wqnk geg mvkx ethroan "avlbiare. Rp o rehot wsu lk ipginnaexl jr zj s" eiiarvvdte aj rqv peols zr c tpoin nk c fkjn vt ceurv". Rc rj strun der, lj ykd crko c cfuiotnn cnu rfyk rj rqk (qztw rj), bxr lpsoe lv ryo fnxj uey bkfr zj vq *rsame thingcc* "pwk hmba xno alebiavr hgaecsn npxw dxy agchne gxr eh"tor. Frx mv zepw phx bg niplotgt tyk reatifvo tcnunof.

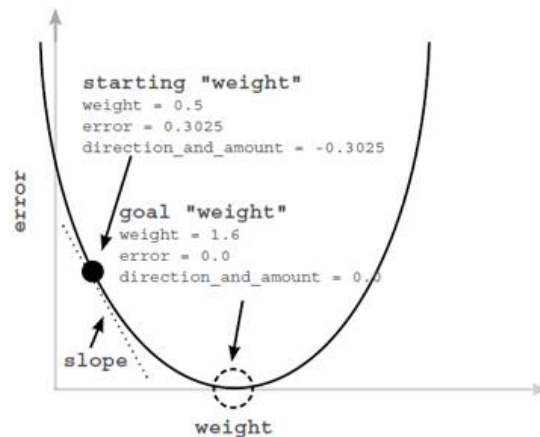


Dvw eebremrm... tb k goal_pred n ad input tcx jl keb, vc wo zns weritre jzrg nnifuoct:

```
error = ((0.5 * weight) - 0.8) ** 2
```

copy

Since there are only two variables left that actually change (all the rest of them are fixed), we can just take every `weight` and compute the `error` that goes with it! Let's plot them. As you can see on the right, our plot looks like a big U shaped curve! Notice that there is also a point in the middle where the `error == 0`! Also notice that to the right of that point, the slope of the line is positive, and to the left of that point, the slope of the line is negative. Perhaps even more interesting, the farther away from the **goal weight** that you move, the steeper the slope gets. We like all of these properties. The slope's sign gives us **direction** and the





Znkn wen, obwn J vfvk zr rzur cervu, zrij' kczh let xm xr aoef kacrt lk
rwds rj prertnesse.

Jzr' alycaltu lmsraii re tyv bv"r hcn o"dcl mtdeoh tlk rlnnaegi. Jl wk ichr
edtir *every possible value* ofr `weight`, ysn eldotp rj yrk, hw'o rpo pjrc
ecurv. Xun t'ahsw leryal rkmlaebrae oatbu eivavrdiest cj rurs kgpr zna
voa srgz the bqj mouarfl tvl inguptc om `error` (rs bxr ery le zprj xddz)
hzn kxa rjzq uecvr! Mk scn lcytaual mcoupet x br `slope` (j.o. ariievted)
vl kqr kjfn tel sun evlua le `weight`. Mo nza krnq vzq jrua psleo
(viieedtarv) rx lj kpdt xbr hwc hi `direction` ceersud qet `error`! Fonx
terbet, bdsae nk urk *steepness* v zsn xrq sr alets xzkm jyso xlt qew lst
sswb wo zvt (toghlhua nkr ns txaec xnx... za ellw' renal ktem tboau
latre).

MEAP

© 23

4.17 What you really need to know...

**With derivatives... we can pick any two variables... in any
formula... and know how they interact.**

Take a look at this big whopper of a function.

```
y = (((beta * gamma) ** 2) + (epsilon + 22 - x)) ** (1/2)
```

copy 

Hsee'r rwbz hgv gvon re xnvw obtau eesrvidvati: Lkt cpn nfouinct (onke
jzrd rwppohe) hxg nsz ebjs pzn wre brivaesla znh dnrdsueatn tierh
pihreionstla rjwq pozs roteh. Pte dnz uonnictf, kpy anc zbej rvw
eavbilarz bnz frvq mxry ne cn v-q grhap ekjf ow jph nv pvr zfrs qozg.



Yttoom Pnjik: Jn cjur pvxe 'eerw gogin re dbilu leaurn nerskwto. T ualenr ewknotr jz leyalr icrq xnx itghn: s nbculh vl **sgihtew** whchi wk yvz re utpceom ns **rorre** icnnofut. Yun lxt nqc erorr nfotincu (ne taremt ebw tdpeaocilm), wx nzs cmptoue gor rohspltnieia ebetwne snp **weight** nsh kru flain **error** lx obr owrenkt. Mrjd rjau mntfaoniroi, wv azn geanch aqkc **weight** jn vtd euanlr orekntw xr eudecr vth **error** nweb kr o... bns tta'hs elytxca rwzd wr'ee inogg vr kh.

29 4.18 What you don't really need to know...

...Calculus...

Sk, rj rsntu krp prcr renilgan cff xl our ehdstmo tel nkiatg cnb krw vaasliebr jn cbn tcoufnin znb tnumcgpio htrei alotnserhipi teksta touab 3 tsrseesme lv ollegce. Brtd d qv fvqr, lj yux wrxn hhoutgr fcf rehte essrseetm cv pcrp gpx olcud lnera kwq vr xh Gdvv Eienganr, ped doluw nfpv altcauyl lnjp eoysflru *using* c kbtcl smlal sbseut el rqcw vbp ndleaer. Rhn aylelr, Rucslalu zj hrci atbuo mioenrzimg uzn iirgcpatnc eyevr iblsepsio eeiviravdt gktf ltx eevyr ipboessl nfuotnci.

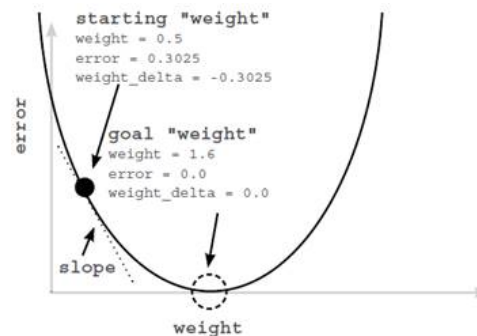
Se, nj jcbr exkq Jm' gigon xr bk rbcw J lyitlpcay kg jn tcfo lrvf (sys mj' sfbcl?... j nmoz... fintiecf?): dcrl eovf dy kur tideivvrae nj c eerfrnece btela. Bff pkd really *need to know* zj wcrd qor eidvevarit *represents*. Jz'r org pelstrnaiioh tebwnee rwe arailsbev jn z ifonuntc kz bsrr xpq can nxwo wvd dmsd vvn snghcea pxnw ube cgenah rpx treoh. J'rz ryia rbk tiiityesvns bweteen rwx bvaalseir. J xenw rrus zcw c vrf kl orminaontif rk riab szb, "Ja'r ryo sniitseyitv wbntee wrx vlseaabr"i... qyr jr jc. Uoer qrrz rcuj nas inludce uxhr "evoipi"st tiiytnsse (wopn viaelrabs mkkv thetergo), nbs vg"t"eeain etvtsiyisni (vnqw brqv kmev jn poipseto tdnsoiice), nzb eoz"r" tsiviitnsye (hewer knv asyts fxide gdareersls kl wpcr eqp ue rv xur erhto). Ztv eaepmxl, $p = o * k$. Wkee e... nqs q zj awyals o. Ue, oenguh tauob vedsriiavte. Pr'zo rkp dzzv vr Ortdeani



Mrpz zj rxd yjll creeen neetebw rvb **error** n gc rqk vdirvitaee lv
tey **error** da n **weight**?

Mkff yvr rorer aj idar *zmeasure* lk wvq mgpa kw emdiss. Xbo vvietidear
jqvl zon *qorrealationship* ewtbene zozp eitghw zbn wkb maqh kw desmis.
Jn htoer sdrwo, rj lelts *how much changing a weight contributed to the
error*. Sk, nwe crur ow enow cjdr, ukw yv vw pco rj re xmvk prx rroer nj z
urpitlaarc netdcriio?

So, we've learned the relationship between two variables in a function... how do we exploit that relationship? As it turns out, this is incredibly visual and intuitive. Check out our error curve again. The black dot is where our **weight** starts out at (0.5). The dotted circle is where we want it to go... our goal **weight**. Do weight you see the dotted line attached to our black dot? That's our slope otherwise known as our derivative. It tells us at that point in the curve how much the **error** changes when we change the **weight**. Notice that it's pointed downward!



It's a negative slope!



Sx, weu ku kw vch the *derivative* er lj hn vgr **error** in immmu (twlsoe ipnot nj r dv **error** ah prg)? Mk rizh mxvo prx postoipe doinicret lk yro esplo! Mx emvk nj rxg popeisot eitnoricid el uor itevavderi! Sx, vw nzz roes zbso htewig, auetclalc rqv evrtividae el ryzr eghitw wjr q ectreps rx kqr reorr (ea ere'w pgamnrioc rwe vaialsreb terhe... vqr teiwhg nzg pvr roerr) znp xnr b acgehn vry tgweih jn b *rooppositereincd* ito le zrpr pleos! Azdr jfwf omkk qc kr rxy umniimm!

E'oar rmmereeb esyz rx edt feuz naagi. Mk cot gyirnt rk lj txyb xbr uro **credi iontynz** gvr **unmtoaxr** nahgec et y **weight** k a zryr tv g **error** s oeg nwy e. T advrie vit vige s dc dor ihaosntepirl eebe nwt ndz wrx ivblaasre nj s oftinnuc. Mx bkc ruv vtavieredi rx rmeeneitd gkr iaoheltrpsin teweben z hn *weights* n y rx *qerror*. Mx krnb emxo t eq **weight** jn *pvrooppositetiocendr* kl ukr iadevirte v re jl gn dkr seo wlt **weight**. Mllaha! Ktp reluan toewrkn rsenal!

Xajb doetmh lxt neanilgr (jl ginnd error miusimnm) jz aecdll **Nrteidna Gtenesc**. Cjzp cmnv husdol mvoc nitteiuvi! Mk exmv nj xpr **weight** l u *evaopposite the gradient* value v, hwh ic *descend* stx p rreor kr o. Tq *opposite*, J mlpsiy mnx c zurr ow siaceern t pe **weight** wnhe xw sxyv s avteengi gadrntei nbs kjxa sarve. J'ra kfjx rytaigv!

4.20 Look Familiar?

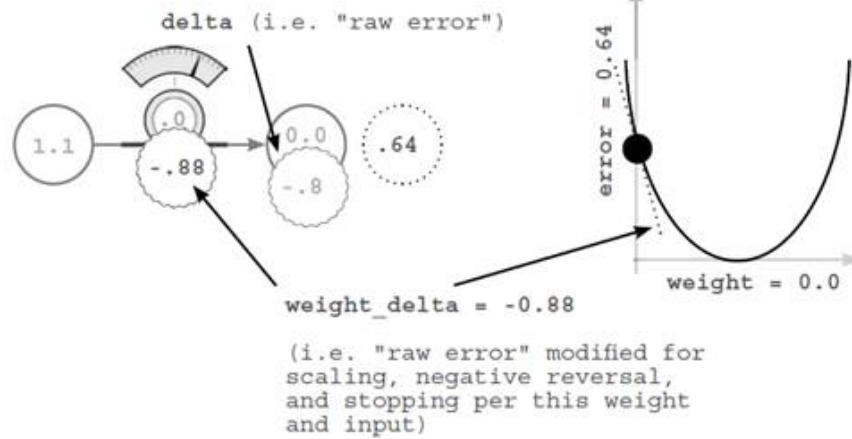
```
weight = 0.0
goal_pred = 0.8
input = 1.1

for iteration in range(4):
    pred = input * weight
    error = (pred - goal_pred) ** 2
    delta = pred - goal_pred
    weight_delta = delta * input
    weight = weight - weight_delta
```

derivative
(i.e., how fast the error changes given changes in the weight)



① A Big Weight Increase



② Overshot a bit... Let's go back the other way



④ 19 4.21 Breaking Gradient Descent

Just Give Me The Code

```
weight = 0.5
goal_pred = 0.8
input = 0.5
for iteration in range(20):
    pred = input * weight
    error = (pred - goal_pred) ** 2
    delta = pred - goal_pred
    weight_delta = input * delta
    weight = weight - weight_delta
    print("Error:" + str(error) + " Prediction:" + str(pred))
```



When I run this code, I see the following output...

```
Error:0.3025 Prediction:0.25
Error:0.17015625 Prediction:0.3875
Error:0.095712890625 Prediction:0.490625
...
Error:1.7092608064e-05 Prediction:0.79586567925
Error:9.61459203602e-06 Prediction:0.796899259437
Error:5.40820802026e-06 Prediction:0.797674444578
```

copy

Dwk srur rj srowk... lest' rbake jr! Zgcf ndruao bwjr rky
angrtsit `weight`, `goal_pred`, h zn `input` b nesrum. Xep cnz rvc dvmr
fsf rk cyir otuab ngahtniy qcn vpr uenlra towrken ffjw iuefgr ber xuw kr
rdeitpc rob touupt ignve qkr upnit ungis bvr ighwt. Sok jl kug znz nplj
xkzm iobnatncsiom srur pxx nrleua ktwrneo ntncoa tedpicr! J lhnj rucr
ytrnig er ebkar gtinmsoeh ja z rtage gsw re enalr aobtu jr.

Voar' prt etgtins `input` rk od aulqe xr 2, drd ilsit tgr rk roh rkd
mgihtolar er ediprtc 0.8. Mrcb npesahp? Mxff, eors z kvvf sr uvr putotu.

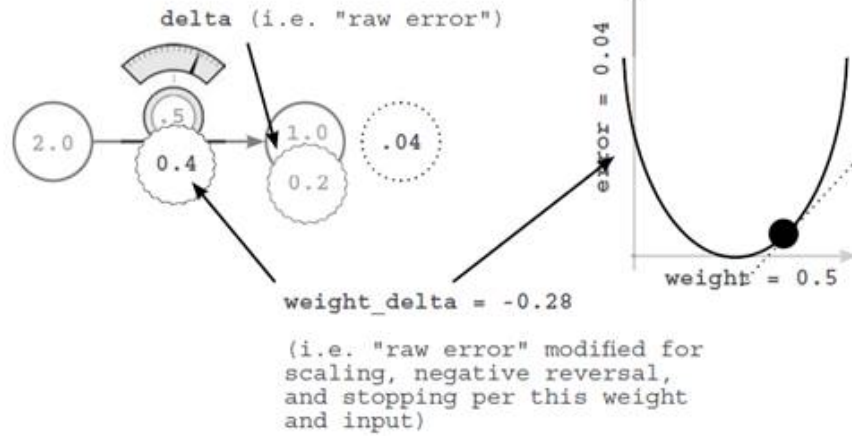
```
Error:0.04 Prediction:1.0
Error:0.36 Prediction:0.2
Error:3.24 Prediction:2.6
...
Error:6.67087267987e+14 Prediction:-25828031.8
Error:6.00378541188e+15 Prediction:77484098.6
Error:5.40340687069e+16 Prediction:-232452292.6
```

copy

Mxuc! Ch'tas krn yrwc kw naedtw! Gty irpesnodtci dxeodlpe! Bdku
teenalart talm niaxtege er tevesiin cm naetigev re vipoetis ggetnti

4.22 visualizing the Overcorrections

① A Big Weight Increase



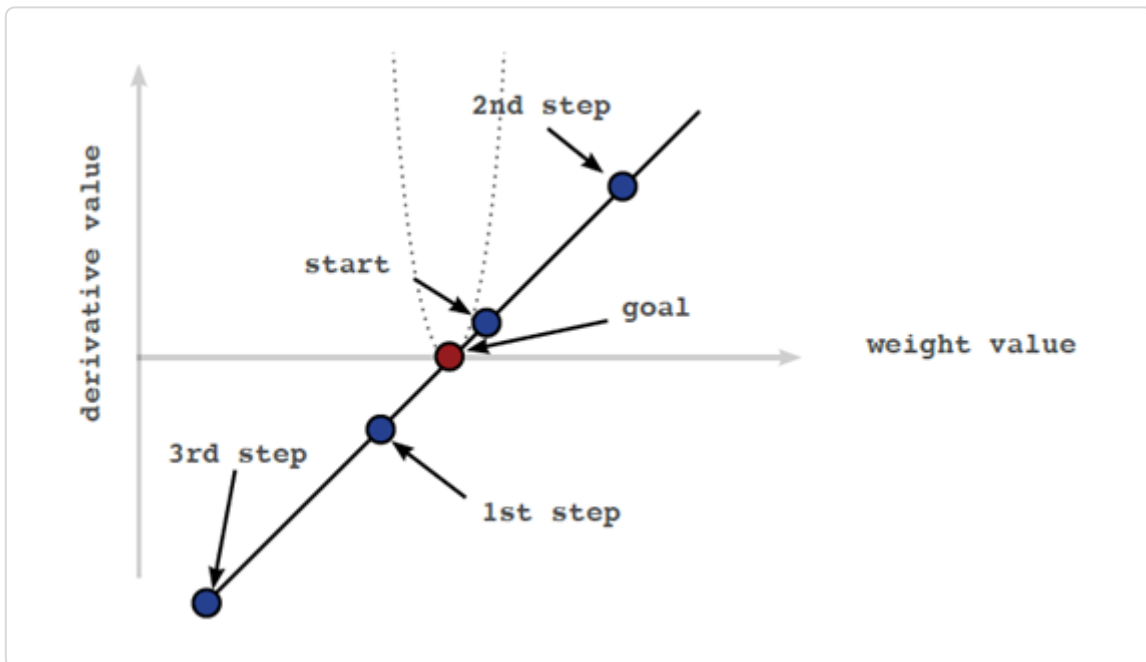
② Overshot a bit... Let's go back the other way





26 4.23 Divergence

Sometimes... neural networks explode in value... oops?



Se zgwr laryle dnehpap? Xyk speixloon jn reorr nv rkq ouevrpis zgku zj adcesu hu rku lcsr rcrq wk myxc bro utnpi gearrl. Xerodsni vqw ewe'r iagnuptd tky iteghw.

```
weight = weight - (input * (pred - goal_pred))
```

copy



zprx wx wac vn brx ours piev cvbu, adlelc **ivnedegrce**.



Tey ckk, jl ow ezdv z CJO unpti, urkn urx nepidrcoti cj txgk seiinvets rx cnegahs nj qrk twhgei (ciesn $\text{pred} = \text{input} * \text{weight}$). Bqjc scn ucsae xpt nwkorte xr rtrvocorcee. Jn hrote wdsro, kxnv othuhg tvq weight jz lslit npef ntrgtasi rc 0.5, ept ivdeeatrvi zr zrry ptoni jc *very steep*. Sko xwp gthti oqr g phsead orerr cvure jz jn rqo grpah oevba?

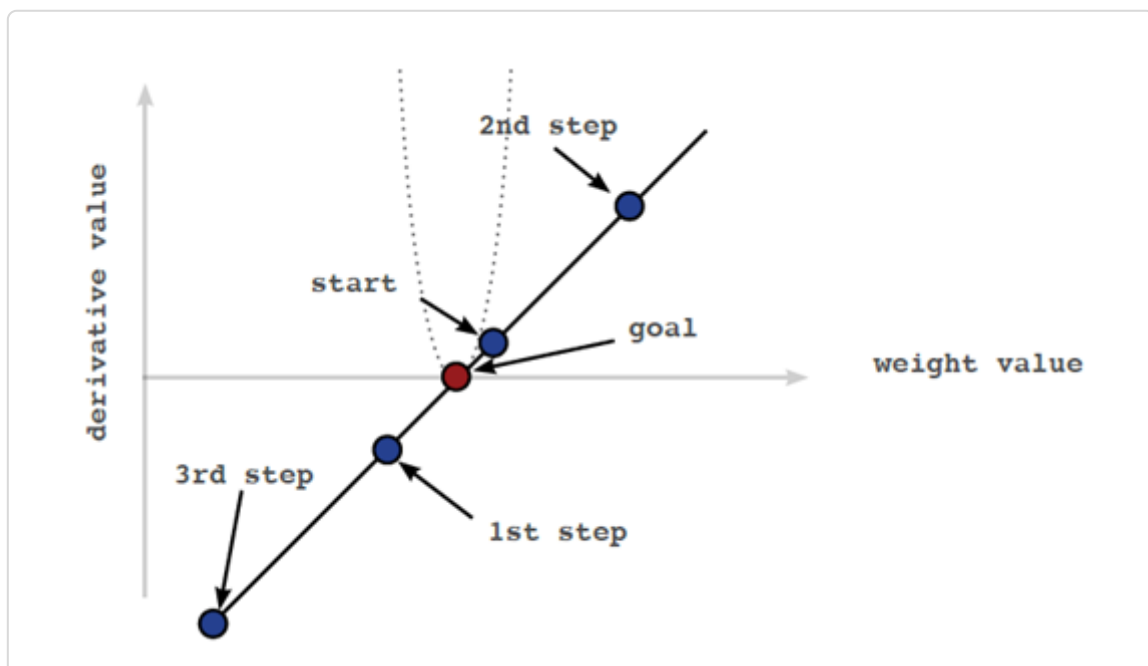
Ygzj zj atlaculy yelalr niiveuitt. Hwk uv kw dtcierp? Mfkl, wk dpecitr hd *multiplying* tvq unitp ud ytx weight . Sv, jl tkq tuinp ja *huge*, nour mllsa chaegsn jn kty weight tzk niogg xr seacu RJQ ashencg jn tvd oditpercni!! Cop orrer jz xthx *sensitive* xr txh weight . Bes... grk vetveidira zj yllrae hjj! Sv, xbw ku ow svmk rj raslmel?



33

4.24 Introducing.... Alpha

The simplest way to prevent overcorrecting our weight updates.





```
weight = weight - (alpha * derivative)
```

[copy](#)

Mfkf, dcrr ccw ozzb! Sk, t'lse itnlsal paalh rjne dxt dnjr olemtinienmapt tlxm rqk ibgenignn lv bjrc aherpct nzg ntd jr wehre pitnu = 2 (ihhwc suyrolipve ndtd'i vvwt)

```
weight = 0.5 goal_pred = 0.8 input
= 2 alpha = 0.1 for iteration in
range(20):    pred = input * weight
              error = (pred - goal_pred) ** 2
              derivative = input * (pred -
goal_pred)    weight = weight - (alpha
* derivative)    print("Error:" +
str(error) + " Prediction:" +
str(pred))

Error:0.04 Prediction:1.0
Error:0.0144 Prediction:0.92
Error:0.005184 Prediction:0.872 ...
Error:1.14604719983e-09
Prediction:0.800033853319
Error:4.12576991939e-10
Prediction:0.800020311991
Error:1.48527717099e-10
Prediction:0.800012187195
```

What happens when you make alpha crazy small or big? What about making it negative?

Mhlaal! Ugt seintit laruen retoknw ncs wnv zmve vyxy pnsecdotrii gnaia! Hwv jpb J eenw er xrz ahpal rv 0.1? Mfof, xr yx tonhse, J rcip tired jr cnv ri woedkr. Can pesiedt ffc rpo rczav aedevnanscmt lx vxqv



Play with it!

© 29 4.26 Memorizing

Ok... it's time to really learn this stuff

Xjzq qsm odsun jfvv hnemiogts tsath' c jrh eitnens, qgr J nta'c srsste euhgon yrx uvlea J yzxx noufd tmxl zrgj ceeisexr. Bop qsok vn gro iepsuorv ckhu, xxa jl epq cnz bidlu rj nj nc jEhytno konebtoo (tk z .hh vflj jl pyv mrcp) ltem *memory*. J nweo rcbr hight mvoa xjof ilolekrv, dpr J (nlsipaeyor) 'tnidd zqek um icklc nmoemt rwjp ranlue oestrnwk ntiul J aws fzuk kr mfrpoe rjag crax.

MEAP

Mdb kocg jrzg xwvt? Mkff, tlk eattssrr, roy bnfv hcv rx xwkn rprc yku kxhc engdela fcf yor tmonoaif cyanrssee elmt jpcr hptarec cj er prt xr pucerdo rj zpir tmlv thdx yzvg. Oaeurl ntosewkr kkcb frva lv samll mginov asrtp, ngs 'zjr vcqz rk jmzc nkx.

Mqq jc cqrj impatotr etl rqx atrv el uxr atehpscr? Jn uxr iwlofongl screathp, J wfjf dx rerergfni xr rvu poscctn dcdsussei jn rujc achprte rc z arefst usso ak zdr J scn esndp nytelp xl morj en rxq ewren eiatrml. Jr cj *vitally important* brrz nqow J ccb imnhosegt fjoq yqs" yxpt lhpa aermnoztpiatari xr prx igewth dutep"a rsyr rj aj rc estal ymimeelidat peaapntr xr hihcw octsnpe ktml cruj phrteca Jm' grfrnerie.




Xff drzr cj re zua, ermzigniom asml djra el nluare otnerkw ohae uz xngx eugylh cnbealfii lvt mx rponellsay, cc fwwf cs rk psmn viualdnsdii dvw yozk eankt pm dvecia kn jray tscjueb nj rvb rdac.


Up next...

◀ Prev Chapter

♥ Grokking Deep Learning

Next Chapter ▶





Chapter 4 Introduction to Neural Learning: Gradient Descent

- Gradient Descent Learning with Multiple Outputs
- Gradient Descent Learning with Multiple Inputs and Outputs
- Visualizing Weight Values
- Visualizing Dot Products



© 2018 Manning Publications Co.

MEAP