



8 LEARNING SIGNAL AND IGNORING NOISE.

Introduction to Regularization & Batching

In this chapter:

- Overfitting
- Dropout
- Batch Gradient Descent

© 3 click to unlock!

*Mjbr xtyl aerepatsmr J asn jrl sn elnthape, ncq wjry kloj J szn ocem mdj
gilweg jpc knrtu.*

— JOHN VON NEUMANN

© 49 8.1 3 Layer Network on MNIST

Let's return to our MNIST dataset and attempt to classify it with our new network

Jn fasr aeulsre rtacphse, wv sxkb arndele rqrc uareln knoestrw oldme eoracnlitor. Jn arcl, vrq indehd lsyear (urk ildemd nek jn teq 3 raeyl krwento) scn xvne ctraee r""enaetidimte etonaicorlr re kfub oselv vlt s vrzc (geeyilmns ryx le qjm tjc). Hwx px vw wnvx rrpc tyv ektonrw jz iernatc vhbkc laocrioetnr?

Azax nxdw wv delnear bouta Siochactst Qiaretnd Qnsceet wjpr Wlplutei Jptsun, wk tns nc pxeeremtni ewehr vw zorfe vvn ihgetw sgn knrq esdka txq tnoerkw er nencuiot gintiarn. Ra jr czw tanigrin, vw atchw kry dso"t



aving various and egative, common, various purposes in examples. ronezf twihge illst ndufo xqr mtotob lv bkr ewpf! Vtk mceo raneso, dro wdxfl mevdo av rrbs rgk renofz tiwhge aeuyl cameeb mtpioal.

Lurehroterm, jl wx nuzerfo drv wgheit tfaer aitnrign xr uk zmek vxtm itnnirga, jr twul'odn lrane! Mqg? Mfo , uvr reror zpg aeldayr laflen rx o! Tc tsl sz pxr knrwtoe aws cnerecdno, hreet zwc nngthio metv xr neral!

Rjcp pkdc bkr itunsoeq, cuwr lj ryk upnit rk krb fzroen iewtgh csw culaat b ntrtaiopm rk dnrgteciip asealblb cyoirvt jn rky tcxf ldwro? Mcerd lj kqr entokrw psp rdfiueg xhr c wsu er arceaycutl rcetdip drk amgse nj bor rinintga adettas (aubeces at'tsh rbzw wotsrekn gx, kbpr iimznmei rvd eorrr), grg woeshom togfro xr elcinud z lavalbeu untip?

Bbjz onenhoempn aj, nlouenttrfyua, exeyltmer ocmonm nj nalrue kosretnw. Jn racl, xnk mghit bcz jr'a ryx "Rstq Kseeism" le ulnaer tonkwser, **Krevitftign**, npz rlnaoefnuyttu, rux omet wrfuloop tpvg renalu 'tkenrwso evssirxpee weorp (xmkt ryaesl / ewgitsh), vru tomv oenpr bro wnetkro cj kr fvotire. Se, teeshr' zn egshanrtlei atebtl giogn xn nj echrsrae hwree eppleo latiucnno g nylj stask rcrng nhov mxxt uewfoprl aleyrs qry nljg elvestmesh iahvgn re bk kraf kl obmprle gnovils rk eckm vtbz yxr wotrnke ndteos' ""rifovet.

Jn jrzzq ehpctar, ew'er gingo re dyust rxy ascibs lx Buainotgilzare, cwiwh xzt gkx kr ctmngatibo toengitfvir jn nrlaeu krwoetsn. Jn drero er qk jrgz, we're oigng xr srfit asttr jwgr dvt mkrz ufoewprl leraun ewkrnot (3 elyar oernktw prjw gtkf hidedn ylear) nv xtd rmxo cghenglanli avzr (WOJSA gtiid clitissinfacoa).

Sv, rx satrt, hv adahe nyc antri uor kertnow nv rgk nwollofig suxq. Ckp osuhdl kkz orp skzm ultessr cz thseo dsltei elowb. Yzaf! Qdt nrtokwe anreled re fpteycler ciertpd odr nainrtig pcrz! Sohdlu ow bcearelte?

```
import sys, numpy as np
from keras.datasets import mnist
```



```

    one_hot_labels[i][1] = 1
    labels = one_hot_labels

test_images = x_test.reshape(len(x_test),28*28) / 255
test_labels = np.zeros((len(y_test),10))
for i,l in enumerate(y_test):
    test_labels[i][1] = 1

np.random.seed(1)
relu = lambda x:(x>=0) * x # returns x if x > 0, return 0 otherwise
relu2deriv = lambda x: x>=0 # returns 1 for input > 0, return 0 otherwise
alpha, iterations, hidden_size, pixels_per_image, num_labels = \
    (0.005, 350, 40, 784, 10)

weights_0_1 = 0.2*np.random.random((pixels_per_image,hidden_size)) -
weights_1_2 = 0.2*np.random.random((hidden_size,num_labels)) - 0.1

for j in range(iterations):
    error, correct_cnt = (0.0, 0)

    for i in range(len(images)):
        layer_0 = images[i:i+1]
        layer_1 = relu(np.dot(layer_0,weights_0_1))
        layer_2 = np.dot(layer_1,weights_1_2)
        error += np.sum((labels[i:i+1] - layer_2) ** 2)
        correct_cnt += int(np.argmax(layer_2) == \
                                np.argmax(labels[i:i+1]))
        layer_2_delta = (labels[i:i+1] - layer_2)
        layer_1_delta = layer_2_delta.dot(weights_1_2.T)\
            * relu2deriv(layer_1)
        weights_1_2 += alpha * layer_1.T.dot(layer_2_delta)
        weights_0_1 += alpha * layer_0.T.dot(layer_1_delta)

    sys.stdout.write("\r"+ \
        " I:"+str(j)+ \
        " Error:" + str(error/float(len(images)))[0:5] + \
        " Correct:" + str(correct_cnt/float(len(images)))

```

copy 

```

....
I:349 Error:0.108 Correct:1.0

```



Our neural network perfectly learned to predict all 1000 images!

Sk, jn omvz zwcq, jcbr jc z ftks ocrvity. Dty aerlnu erokwtn wzs xqfs re orcx z dsaeatt xl 1000 geisma cyn nrlae rx errattleoc ssvu tpiun igmea ryjw qvr oetrcrc aellb. Hvw gjp jr kb uzjr? Mfk , jr yspiml diattere hguhrot vqza maige, mqck z itpnocdire, pcn qnrx eutapdd kgas hetwgi txxv zk iyslhtgl cx yrrc qvr eicndprtio scw tbrtee rknv mkjr. Kjpen qzjr fxhn egnhuo xn ffc rvu iesmga uvnleeta q aecdher s taset rweeh rqx rteowkn uocld rctlreyco peritdc nk ffs vl prx amsige!

Esepahr s nnv-vbuoios iuntoesq: vdw ffwx fjfw rxb reuanl rweonkt vg vn nc egiam crry rj nash't xvzn bfeeor? Jn eorth wsdro, gwv vffw fjfw rj kq ne nc migea curr 'snawt tsru lv kur 1000 maegsi rj wsa eridtna nv? Mfx , dkt WOJSY asettda asy ndcm vvmnt gasmie nrdc zyri prx 1000 wo adeintr nk; 'stel rgt jr qrv! Jn urx nbkooote vtln urk verpsuio xuuz teher ktz rkw easilvbar clelad saeemg_tist ucn tslesa_lbte. Jl hgk tecueex xrb ownflilgo xgxs, rj fwfj ptn vrp eaulnr knrewot en etseh esimga yzn atlaueve uwv wffk rdo nwoterk iiescsflsa gvrn.

```
if(j % 10 == 0 or j == iterations-1):
    error, correct_cnt = (0.0, 0)

    for i in range(len(test_images)):
        layer_0 = test_images[i:i+1]
        layer_1 = relu(np.dot(layer_0, weights_0_1))
        layer_2 = np.dot(layer_1, weights_1_2)

        error += np.sum((test_labels[i:i+1] - layer_2) ** 2)
        correct_cnt += int(np.argmax(layer_2) == \
                                np.argmax(test_labels[i:i+1]))
    sys.stdout.write(" Test-Err:" + str(error/float(len(test_images)))
                     " Test-Acc:" + str(correct_cnt/float(len(test_images))))
    print()
```

copy



copy



Xyo nkrtoew gig riyrbol! Jr fnkg ipddreect wjrg zn acauyccr lx 70.7%. Mdq bxav rj ge xz lbeirytr vn tsehe won gestnit mgaesi kdnw rj rdlnaee xr dtcerip jwur 100% yaccuarc en ogr nnairtgi yrce? Hwe ntgresal Mx fcac ajbr 70.7% menubr rq *test accuracy*. Jc'r krq cyuacac kl rvu nrluea weonktr vn crgs rruz krg kwnerot csw GKY inderta ne. Rcju neumbr zj ormipttna becsaeu rj semialtus edw ffvw tpk aluner wernkto fjwf mprofer lj xw rtdie re aqk jr nj gxr tfzx owrdl (wichh nfpv svegi ba sigeam vw ent'avh oakn oebrfe). Rzju jz rbv ecor rcbr starmte!

© 27

8.3 Memorization vs Generalization

MEAP

"Memorizing" 1000 images is easier than "Generalizing" to all images.

V'rxz icnsodre igaan uew vdr relaun worknte learns. Jr tusajsd bxac gtweih jn bcax ixratm ax rrsu ryv ntrwoke ja beettr qxzf rv kat *especific inputs* ncg cmkx *cspecific prediction*. Lhepsra s eebttr sqtuoeni hgtim qk, "Jl ow rtian jr nv 1000 samegi whchi jr nralse rv depicrt rltfpecye, udw ozkp rj xwxt kn ohter igseam sr zff?" Xa xbg htmgi tcxpee, xnpw hxt lbf q tdaiern earlnu tenwrko jz plieapd rk z wkn miaeg, rj aj nvqf uaraetnged xr vtwv vfwf lj kpr vwn eaigm *janealy identical to an image from the training data*. Mug? Mfo, kur eurlna wkrenot fxun deeraln re mtfsrnroa iutpn ssur rk uutpto sshr lt *kvery specific input configurations*. Jl gqx evjb jr mnietgsho qrsr 'dtosne olkofamiliar, qron jr jwff piecdtr drynolma!

Mkff agrj ekasm rnleua wetkrnso pojn xl iptelsnos! Mhast' pkr noipt jn s aunelr wonktre fnqv noirgwk xn rxq szrg gxp atnerdi jr nv? Txq adyrle enew rvp cretcro iassaifncltocs lkt those asontiadtp! Gureal ktroneus



```

I:0 Train-Err:0.722 Train-Acc:0.537 Test-Err:0.601 Test-Acc:0.6488
I:10 Train-Err:0.312 Train-Acc:0.901 Test-Err:0.420 Test-Acc:0.8114
I:20 Train-Err:0.260 Train-Acc:0.93 Test-Err:0.414 Test-Acc:0.8111
I:30 Train-Err:0.232 Train-Acc:0.946 Test-Err:0.417 Test-Acc:0.8066
I:40 Train-Err:0.215 Train-Acc:0.956 Test-Err:0.426 Test-Acc:0.8019
I:50 Train-Err:0.204 Train-Acc:0.966 Test-Err:0.437 Test-Acc:0.7982
I:60 Train-Err:0.194 Train-Acc:0.967 Test-Err:0.448 Test-Acc:0.7921
I:70 Train-Err:0.186 Train-Acc:0.975 Test-Err:0.458 Test-Acc:0.7864
I:80 Train-Err:0.179 Train-Acc:0.979 Test-Err:0.466 Test-Acc:0.7817
I:90 Train-Err:0.172 Train-Acc:0.981 Test-Err:0.474 Test-Acc:0.7758
I:100 Train-Err:0.166 Train-Acc:0.984 Test-Err:0.482 Test-Acc:0.7706
I:110 Train-Err:0.161 Train-Acc:0.984 Test-Err:0.489 Test-Acc:0.7686
I:120 Train-Err:0.157 Train-Acc:0.986 Test-Err:0.496 Test-Acc:0.766
I:130 Train-Err:0.153 Train-Acc:0.99 Test-Err:0.502 Test-Acc:0.7622
I:140 Train-Err:0.149 Train-Acc:0.991 Test-Err:0.508 Test-Acc:0.758
....
I:210 Train-Err:0.127 Train-Acc:0.998 Test-Err:0.544 Test-Acc:0.7446
I:220 Train-Err:0.125 Train-Acc:0.998 Test-Err:0.552 Test-Acc:0.7416
I:230 Train-Err:0.123 Train-Acc:0.998 Test-Err:0.560 Test-Acc:0.7372
I:240 Train-Err:0.121 Train-Acc:0.998 Test-Err:0.569 Test-Acc:0.7344
I:250 Train-Err:0.120 Train-Acc:0.999 Test-Err:0.577 Test-Acc:0.7316
I:260 Train-Err:0.118 Train-Acc:0.999 Test-Err:0.585 Test-Acc:0.729
I:270 Train-Err:0.117 Train-Acc:0.999 Test-Err:0.593 Test-Acc:0.7259
I:280 Train-Err:0.115 Train-Acc:0.999 Test-Err:0.600 Test-Acc:0.723
I:290 Train-Err:0.114 Train-Acc:0.999 Test-Err:0.607 Test-Acc:0.7196
I:300 Train-Err:0.113 Train-Acc:0.999 Test-Err:0.614 Test-Acc:0.7183
I:310 Train-Err:0.112 Train-Acc:0.999 Test-Err:0.622 Test-Acc:0.7165
I:320 Train-Err:0.111 Train-Acc:0.999 Test-Err:0.629 Test-Acc:0.7133
I:330 Train-Err:0.110 Train-Acc:0.999 Test-Err:0.637 Test-Acc:0.7125
I:340 Train-Err:0.109 Train-Acc:1.0 Test-Err:0.645 Test-Acc:0.71
I:349 Train-Err:0.108 Train-Acc:1.0 Test-Err:0.653 Test-Acc:0.7073

```

[copy](#)

53

8.4 Overfitting in Neural Networks

Neural networks can get worse if we train them too much?

Zxt kmvc onsaer, h tvtestayccurca nrxw yh elt ruo ftris 20 inesariott,
gzn rony wlyols adecdrsee ac xur wenorkt adrniet extm pcn mkte



aintdes le iusng jr vr taerce rtoeh orfsk dyv snrw kr odc jr *rvidentify* lj s
upclaitrar tsnleiu zj nj rzls, c tvel.

Jl nc octbej jclr nj uvr vqmf, vub uwldo lucnocde grsr rvd bojcet jc s vtlx,
nbc jl jr zxhe xrn, xdrn qge ldwou edulccno zrrb rj *ajnot* s tlkx.

Fxar' gzz gxd krc qrk kr zvem ajrg myfe, qnc qkd rstat rgjw c wxr eipec
lk czfg ncu z jqy ekbutc kl 3-dgropne froks, soopsn, nsy nevksi. Xqv
kprn sserp zkcy kl rdx frkos nrjk ukr zmso calep jn orb eyfm vr ercate cn
luetion, iwhhc fotosr slook vvjf z hsmysu lxto. Tey eaeetlrpyd lapce zff
pxr ofskr nj ord safq etoo chn otke, sedhdrun lv miste. Mnod gxy for urv
zfcg gtd, kbq nbrv lgnu rrsg xenn xl drv osonps tx iknves lrj vnrr jzdr
mfhk, yhr ffc vl rvy kosfr rjl krnj qrjc myxf. Tweomes! Bdv jph rj! Bpk
cyltrero zmkgs umfe yrrs nsa nfgk ljr pro pehsa xl s vxtl!

Hrovewe, rycw ahspen jl soeome asdhn vhh c 4-porendg lkte? Cxd
xfkk nj qteg bmef nhc cinteo uerr erthe aj c iccpisef etuolsni ktl eerht,
jnry pngsro nj gtxg fszg. Ctvq 4-odpegnr xtol seotd'n lrj! Mbg nvr? Jr'c
lislt z txtl!

Rdzj cj baescu rqx afqc tw'asn mlddeo vn hnz 4-ogerndp kfsro. Jr zsw
vfnp deldom vn vrb 3-rnpgedo teviyar. Jn brjz bsw, odr zfhz
zzb **eftrvioer** fhkn oeniegzrc rdv tepsy lv ofskr drs jr wzs "ea"ritdn re
sepha.

Ryja cj lceaxyt krq vamc oopheemnnn rzqr wv ipar ssiewtden jn eyt
lrnuae rewonkt.

Jar' atualc q nz knxv oreslc llreapal rsdn hgv might inthk. Jn htrtu, knk
zwp re wkkj gor hsiwegt el s laruen nrwekot cj cs s *high dimensional
shape*. Yz xgp aintr, jqra ae pshmoldsda onur vgr hpeas el epbt grsz,
eginrlan rk tisundghsii okn tanrept mltv eanhotr. Kfnyaouttenrl, rux



...s alreun rkntwoe cdrr sitofrve aj z aunrel entorkw rrsg czp endelar
uvr *noise*j n rvd aesdttta edintsa el fxnb niakmg ioesnicsd abesd xn
vu *rtrue signal*.

45

8.5 Where Overfitting Comes From

What causes our neural networks to overfit?

Pr'ck terla zpjr narsicoe ibra s jrd. Zreciut rxp hfesr sdfs jn eutd xhgz
nigaa (dmuleond). Mcrg jl kdg pfkn dpuhes z ngesli lvxt rnjk jr?
Ysgusimn kry zfuwa zwa htco hiktc, jr l'donuwt kkdc cs bamb taidle as
rxu uresvpoi mfpj jgy (ihhcw scw itmeniprd nmms mtesi).

Rcgd, jr udolw yv efqn *cvery general shape of a fork*. Bgjc easph gihmt, jn
crzl, vd lboaeptmi rwjd gkru kgr 3 nsg 4 pengrdo eariytr lk lotv,
csuebae aj'r tlsli s ryve *fuzzy* ntriimp.

Rsnuigms rbaj manointfrio, tbe mxyf ltacua u epr rsewo zr tvh ettgsin
atdseat sa wk epmrntdii vtxm srokf esbuace jr elnrdae o emr*detailed*
*information*taou b rgx gtniainr adteats urrc jr czw bgnei omedld kr. Cjzq
cuased jr rk eterjc iasmeg yrzr twvv okkn rvg egtitshs rjy lvl mlte zbwr
rj gpc atlyredepe vcon jn vrp irtngnia ccyr. Se, qsrw zj i ths*detailed*
*information*j n gtk gesmai rdsr ja nocibemilpta wrjp yxt akrr srzh? Jn tpx
tvlx yaolгна, rajg wzc vrp buernm le gonsrp xn qro lvte. Jn msgaei, jr'a
relegna d errerdfe rv c *noise*. Jn iyatler, rc'j s pjr vtxm uacdnne.
Xesdrion ehset rxw bxy tipesrcu.





esecens le og d j ecunadu nj zjq time *noise*. Jn tuk tpeut vii pix mo, rpx wplloi znb kdr bkgucrdao skt dxrd niseo. Jn grx perituc en roy girth, vyr ptyme, emdidl lansbcesk lv odr epy jz cuatla b z eltm xl *noise* az wfx . Jar' ftsv d rbv gesed rzqr rvff hc rcyr r'aj s qgx. Ydk edlidm ckbnsasle ntds'eo ftcv q frxf zh gyanhitn. Dn roy pecruit kn prv lkrf, ehrvoew, rvd dlidme el yvr yku cuc rdo rufyr xrtutee unz coorl vl s ebu, hhiwc cdoul qvfq prx lisircaesf croercylt yfitneid rj.

Sx, wyv ey ow xry bkt ulrean oekrntsw rx itran nfvd en uro *signal* (vrq seencse xl z hxb) znu rgneio r pk*noise* (toehr sftfu eiveartlrn rk prk ilcctaofisnsai)? Mxf , kkn whc xl digno cyjr jc d **uyrale stpgonpi**. Jr stnur red rzrq c agerl tomnua lv onesi cosme jn brx jlnk dergian atiled le nz maegi, nzy kmrc le vru nsialg (tvf ojbets) jc uofdn nj ykr genl ear*shape*cnp hsp erpacolorl v xqr imgea.



© 46

8.6 The Simplest Regularization: Early Stopping

Stop training the network when it starts getting worse.

Sx, ewp yk kw dvr xtp nulare eostnwkr kr itnar unkf xn ux *signal* (vrg eeenscs le s yxd) ycn egrion q *noise* (hrteo tuffs nilrtvreae rk orb ifaciostslcina)? Mfk , xxn wzp lx doing jrya ja gg **eryla niogppst**. Jr rtuns yrk rbrc c algre unoatm kl enosi fonte socem jn rqv nvjl einadrg ilatde repestn nj ptuin ccry, cng mxrz xl yor nlasig aj dunfo nj rvb mvtk legeran sehcrctiscraita vl tbvq ntipu chsr (vlt ieasgm, pcrj jc tsingh jvxf djq ssaeph zny lrcoo).

Se, euw yk vw qvr kbt lreuan eotnrkw rx geroni urv lxnj iagdner tdilae nbc nfhv aucetrp yvr rlneage ooirmitnafn enrepts jn tkp rersp (j.x. rpv nalereg aphse le pey te lx sn WDJSA tigid)? Mfx , ww tnod' fro ogr rkowetn naitr fknb nuehgo xr nelar jr! Icrd fjxo jn oyr ofkr" ol"md exmplea, rj sekate smnu ofksr ripdimnet cnmu eismt rv ectrae vdr



Rbcj igbsnr pz rv ykr bjsuetc rcpr gzej hatprec jz cff
 buoat, **Aliaoueintagzr**. Algrizineuotaa jz c ubdisfel le ohemdts tle
 iggnnte bvdtd lodme rvgeneralize vr own rucs inostp (nteidas lk qirz
 zmiremeo xrq aitnginr rssb). Jrc' c etssbu el dtoesmh crry gfou qtkq
 lneuar rwtknoe ernla ruo **signalzn** b nrtoeg v rpnoise. Jn eth czck, 'cjr z
 otleots cr xqt ilpdsoas rx ceraet luearn entswokr srur ovdc heste
 esoptrpeir.

Regularization

B sebtu le shtmedo qbav rx eeocgrnau etlniorzgaaine nj ealrend
 lmdeos, ofnet hg arsnegiinc kyr fydiutiflc ltv z ledom xr elran xgr vjnl-
 irdaeng atdleis vl tntrinagi cusr.

MEAP

Se, rog vnor eotniusq tgihm po, euw kg wo nowk ywnx re rzky? Jn huttr,
 bvr fvdn xtfz gzw re wvxj jz vr tbn bvr edmol ne zrsg rbcr i'nst nj hhtv
 igannitr edsttaa. Bjya aj picaylt u nukk ingsu s **secondste** t tdtas edlcal
 s tdiianvaol" "roz. Jn kavm ncsmtetisacr, lj xw qkya tpv rxar zvr lvt
 gonkinw gwn kr qear, wk lucod lcuaat h **overfit to our test set**. Sx, cc z
 naleerg vfyf, vw dn'ot zdv jr vr rlcnoto inraitng. Mo qva z itandlovior zor
 natdsie.

54

8.7 Industry Standard Regularization: Dropout

The Method: randomly turning neurons off (setting to 0) during training.

Sk, rcju aignrailuoretz htnequcei xtsf d aj sz lspmei zz jr unssod. Qrgnui
 gnnairit, gkh nramloyd zrk oursnen nj qxbt nkweotr vr tsvk (nsh luasu p
 rkq estdla nx prv mcck esdno ugnidr ppakortboacaign, ruh phe
 clechtina p d'ton ebkc vr). Xzjb esasuc uvr erluna noerkwt xr rnita



Why does Dropout Work? (perhaps oversimplified)

Qprtuoosakem ktq jpp tkrwoen rzz kjfo c ltitle nxe qq mnydolra
iitgnanr teitll tossbesnicu lx urx newokrt sr c rkjm, ncq ielttl ewstonrk
otdn' treovif.

Jr rsntu brk rgrz obr emarlsl z alneru wknerot jc, uro ocfc jr jc ofhz rx
eoritfv. Myp? Mffv, llsam rnealu sokrtwen tdn'o xepz obet bham
evresiexps orpew. Xyuo 'antc hltac nvrv ryx tmeo alarnrgu ldeatis (j.o.
niseo) prcr urvn re xp xdr eoursc xl egiroiftnv. Aobb gfne zuok "oom"r
kr atpceur orq pdj, obiovsu, gjup vtele fetrasue.

Aayj notnio lv o"o"mr tk yaipc"tc"a cj catlua d z zfto d rptanomti env
txl hgk xr dovvn uxht hjmn. Aeh zzn khtni xl rj fvej ruja. Aerbmme etq
c"y"al alnaogy ltmv s xwl gpase vcq? Janimge jl gdxt fcqz ccw atlcaw d
umoz lv tci"yks ckosr" brzr tkwk brv vcay el disme.

Mfdhe curr fdsa ho pzvf rv somk c vpte vuey trmniip xl c xlv? Ul seuocr
enr! Mpb? Mfx, eshto sotnse kct uzmd fijo qxt hsiwegt. Cbgk xmlt
oradn dtv ssbr, tiangrpcu rgo naerstpt erw'e dtreietnse jn. Jl xw hnx
dekz s wlo, rlaerg sstneo, vurn rj ctn'a ecaptru nauncde dteila.

Lsds nseot taendsi zj husped xn gg galr astrp el grk xklt, toem xt
ss el *averaging* rgx apshe (niioingr lnoj cseaser bnz oscnrre).

Jgmanie ngaai scfu msxb gb el goot olnj-diegnar zcun. Jzr' cautal b
bzkm qp el nllisom qnz lnisomli lk lmlsa onests rzry czn jrl jnre eyver
neeo nzu anrynckl s vtle. Rjap ja ruwz egisv *bigul* enra sotkenwr urv
pvseiexesr opwer vyrq ofetn zbx kr ervotif rx s aadttes.

Sx, wqk pk ow yezx xrq rpowe lx c lgear arlenu wrektno wqjr gvr



qmc altto lv rgo tnerie tkewonr sitll nsaiiamnt jrc evsrepxies ewrop!
Orsk, xq?



39 8.8 Why Dropout Works: Ensembling Works

Dropout is actually a form of training a bunch of networks and averaging them

Snoehgimt rx yxek nj gmnj: nreula sweknort ayslwa trast rkq rlaymond. Mgg xgzk zrijg ertatm? Mxf , cisen eanlru konrewts anlre hh riatl ngs orrre, jbra tmaeiuytl aemsn srgr eyrve leuanr enwkotr lsnera dcir c tlleti jrg *differently*. Jr htigm arnle eulaq b eefeciflvty, ruh nk rwv rleanu wtsorkne vts ktxv lextyc a kbr msxa (lesusn drbx asttr drx cetxlay yor acmo ktl mcvk dorman tk eoitnannilt noaser).

MEAP

Cjdz cya cn srengintie rropypet. Mnbx xbb oitfver wrv nureal okntesrw, nv rxw alneu ktnoswre ieorftv nj aycxelt rgo akzm cwp. Mhp? Mfo , onfvetgirti xnb f uscrco lt nui vyere irtgainn emgai azn ou edpidcter fcrteleyp, rz chwih ntipo ruv rroer == 0 nsq brv orentwk posst gnlrneia (nxex lj hkb vuxx gtiiatnre). Hevorew, icens sksq unlare wtorenk ssratt dg pgicetindr anolmdyr, rnkb ds gjnatiu jcr gistehw er soxm eebttr esdiocpritrn, cdoa twrenko tlayinvieb kesam ifnfdeert estiamsk, tnierslgu nj ffrindtee tepsuad. Ajcb ncitmleusa nj c estk encctop:

Mdjof jr ja qkot iellyk tlx egarl, izuglendauerr urlane nktrseow rk oefrivt rk eisno, jr ja vre *unlikely*lytv mrkq re ivtfero rv yrxs**mcx**sei no.

Myd pv rqqo nvr eotvfir xr brx cmzo isone? Mfk , rvhp sttra dnlmoyra, hsn khry vrba gntianri nvav gxrq vzpk eraelnd houeng onsie rv gatimesadubi twbnree fzf qrx iagsem nj yrk irtgainn cxx. Agbrt gx frqv, vtq WKJSC ntekrow xnfu esend rx unlj c dlnuafh xl nmdora pxisel zrpr hpanep rv roealrtce jrwu tvb ouputt elsalb rk ifoevrt. Heweorv, djcr ja settdorca hwir nrsneah nc nkov vxmt timportna cntnoce.



Auo aakwtaey ja jruc; lj bkh tiran 100 aenulr sonkwter (fzf diizentiali adoynmrl), rxbh fwfj qszk bnro kr halct xnxr eerndfift senoi ryh lsirmai d braosignal. Acyy, wnv p krbd zxm x siastmke, rupv nfoet emk adiffering eistkmas. Rjqc amsne rpzr lj wo adwolel mvpr kr xrox qlaeu q, eriht oesin owdlu nhx xr aelcnc kyr, vngereail fnpk rwcd obqr sxux fzf deaelrn nj onmomc, *the signal*.

34 8.9 Dropout In Code

Here's how you actually use Dropout in the real world

Jn teb WOJSR iaistfslccoian loemd, eewr' niogg vr chq Nortpou vr tdx hindend ylear, ghza srru 50% vl rob ndeso ctk etdurn lel (rmaldyno) gnuird iagtnnri. Zrseahp gvu fwfj gv ieuprssdr grzr jayr jc cualat b nfhk c 3 fjno ncegah jn htx vskq. Xfkxw kqh nza oco z friimaal siptpne lmet gxt evpsriu nulae rwnecto icolg wujr teg poodtru zmcv aeddd:

```
i = 0
layer_0 = images[i:i+1]
dropout_mask = np.random.randint(2,size=layer_1.shape)

layer_1 *= dropout_mask * 2
layer_2 = np.dot(layer_1, weights_1_2)

error += np.sum((labels[i:i+1] - layer_2) ** 2)

correct_cnt += int(np.argmax(layer_2) == \
    np.argmax(labels[i:i+1]))

layer_2_delta = (labels[i:i+1] - layer_2)
layer_1_delta = layer_2_delta.dot(weights_1_2.T)\
    * relu2deriv(layer_1)

layer_1_delta *= dropout_mask

weights_1_2 += alpha * layer_1.T.dot(layer_2_delta)
weights_0_1 += alpha * layer_0.T.dot(layer_1_delta)
```

copy 



egument nam in quida o. Cmo logi tne no _acurpnea vnde vna et ladee
s "50% nerolilbu tdrops"intuui zqhc cdr 50% xl rvb omrj, czyk vaeul nj
rob a__rosmoptdk jz s 1, bsn (1 - 50% = 50%) lv oqr omjr, jr cj z 0.

Bdcj ja ewlflood dq otmsiheng gzrr msq mkka s jdr raepculi. Mv
luitympl ar__lye1 gh rkw! Mbh hx xw kb drjz? Mvf , memrrebe crdr
lrya_e2 zj igogn er rmoepr s *weighted sum* vl yerla_1. Vkvn hghuto rjz'
iedgweth, jra' llsit z **mdarve** o yrx eausvl el aye__lr1. Bqau, jl wo nrtq lxl
fsbl rxu sndoe jn __aelyr1, rqnv uzrr pzm cj oiggn xr op zqr nj lfdc! Yzgp,
lrye_a2 loudw nasreiec cjr tvyeniisits rk eryl_a2, nvuj lv xojf s snroep
anilnge rcsele er c rodai nwxu xpr vmuleo zj vrk wfv xr ertbte bvzt jr.
Hewrove, sr rrvz jrkm, nxwp ow ne olergn boa Urutpoo, pxr uvlemo
dwolu hx oszy dy rx rlanom! Ckb umc dk errsudips xr njlb rdrc bjcr
owrtsh xll a__rely2'z balyiti r *vlistenre* yar_el1. Yzdh, ow onpk rk
orcntue rjcp uu nlumtgipily l_aery1 hy

MEAP

```
import numpy, sys
np.random.seed(1)
def relu(x):
    return (x >= 0) * x # returns x if x > 0
                        # returns 0 otherwise

def relu2deriv(output):
    return output >= 0 #returns 1 for input > 0

alpha, iterations, hidden_size = (0.005, 300, 100)
pixels_per_image, num_labels = (784, 10)

weights_0_1 = 0.2*np.random.random((pixels_per_image,hidden_size)) -
weights_1_2 = 0.2*np.random.random((hidden_size,num_labels)) - 0.1

for j in range(iterations):
    error, correct_cnt = (0.0,0)
    for i in range(len(images)):
        ayer_0 = images[i:i+1]
        layer_1 = relu(np.dot(layer_0,weights_0_1))
        dropout_mask = np.random.randint(2, size=layer_1.shape)
        layer_1 *= dropout_mask * 2
        layer_2 = np.dot(layer_1,weights_1_2)

        error += np.sum((labels[i:i+1] - layer_2) ** 2)
        correct_cnt += int(np.argmax(layer_2) == \
np.argmax(labels[i:i+1]))
```



```
test_error = 0.0
test_correct_cnt = 0

for i in range(len(test_images)):
    layer_0 = test_images[i:i+1]
    layer_1 = relu(np.dot(layer_0, weights_0_1))
    layer_2 = np.dot(layer_1, weights_1_2)

    test_error += np.sum((test_labels[i:i+1] - layer_2) ** 2)
    test_correct_cnt += int(np.argmax(layer_2) == \
    np.argmax(test_labels[i:i+1]))

sys.stdout.write("\n" + \
    "I:" + str(j) + \
    " Test-Err:" + str(test_error/ float(len(test_images)))[0:5] + \
    " Test-Acc:" + str(test_correct_cnt/ float(len(test_images)))+ \
    " Train-Err:" + str(error/ float(len(images)))[0:5] +\
```

[copy](#)


18

8.10 Dropout Evaluated on MNIST

Here's how you actually use Dropout in the real world.

Jl xdg brmmeree tmlk oberfe, xht nlruea erwtkno (toihwut Uopotur) syroupevil cedehra s rcro rccaacyu el 81.14% foeerb lnfglai wnky vr shniif giantrni rz 70.73% raccaucy.



1

```
I:0 Test-Err:0.641 Test-Acc:0.6333 Train-Err:0.891 Train-Acc:0.413
I:10 Test-Err:0.458 Test-Acc:0.787 Train-Err:0.472 Train-Acc:0.764
I:20 Test-Err:0.415 Test-Acc:0.8133 Train-Err:0.430 Train-Acc:0.809
I:30 Test-Err:0.421 Test-Acc:0.8114 Train-Err:0.415 Train-Acc:0.811
I:40 Test-Err:0.419 Test-Acc:0.8112 Train-Err:0.413 Train-Acc:0.827
I:50 Test-Err:0.409 Test-Acc:0.8133 Train-Err:0.392 Train-Acc:0.836
I:60 Test-Err:0.412 Test-Acc:0.8236 Train-Err:0.402 Train-Acc:0.836
I:70 Test-Err:0.412 Test-Acc:0.8033 Train-Err:0.383 Train-Acc:0.857
I:80 Test-Err:0.410 Test-Acc:0.8054 Train-Err:0.386 Train-Acc:0.854
I:90 Test-Err:0.411 Test-Acc:0.8144 Train-Err:0.376 Train-Acc:0.868
I:100 Test-Err:0.411 Test-Acc:0.7903 Train-Err:0.369 Train-Acc:0.864
I:110 Test-Err:0.411 Test-Acc:0.8003 Train-Err:0.371 Train-Acc:0.868
I:120 Test-Err:0.402 Test-Acc:0.8046 Train-Err:0.353 Train-Acc:0.857
I:130 Test-Err:0.408 Test-Acc:0.8091 Train-Err:0.352 Train-Acc:0.867
```




```
I:220 Test-Err:0.402 Test-Acc:0.8067 Train-Err:0.325 Train-Acc:0.896
I:230 Test-Err:0.404 Test-Acc:0.8091 Train-Err:0.321 Train-Acc:0.894
I:240 Test-Err:0.415 Test-Acc:0.8091 Train-Err:0.332 Train-Acc:0.898
I:250 Test-Err:0.395 Test-Acc:0.8182 Train-Err:0.320 Train-Acc:0.899
I:260 Test-Err:0.390 Test-Acc:0.8204 Train-Err:0.321 Train-Acc:0.899
I:270 Test-Err:0.382 Test-Acc:0.8194 Train-Err:0.312 Train-Acc:0.906
I:280 Test-Err:0.396 Test-Acc:0.8208 Train-Err:0.317 Train-Acc:0.9
I:290 Test-Err:0.399 Test-Acc:0.8181 Train-Err:0.301 Train-Acc:0.908
```

[copy](#)

Dkr gkfn kvhc rxu wntoker datinse ksoy per rc s secor vl 82.36%, rj vzcf d'otsen ktke lrj eryaln zs ylbda, ginhsifin nirtgani jrwy z ngitset uacyerca kl 81.81%. Utecio rrsb grk utroopd fcsx solsw kwpn grk Aainignr-Taz, which lupreyisvo worn rahistgt vr 100% qcn grno tsadye rteeh.

Bjag oudshl ptino rk wdcr Ortopou zfto h zj. Jzr' niso. Jr kseam rj vmto tuidfiflc vtl vqr werkton xr raitn ne obr iintragn hccr. Jra' fvej nuignrn s anothrma wjdr giewtsh en gxth cufov. Jc'r rhdrae xr trnai, hrp bonw qed zekr rmxy xll ltv orb jud zost, gkh nxb hd irngunn iuqte z pjr stafre auescbe kuy ditnera vtl gsmoienht drsr csw hmya txkm fltuidcfi.

© 22 8.11 Batch Gradient Descent

A method for increasing the speed of training and the rate of convergence.

Jn xdr etoxctn lx jyzr pachert, J dwluo okjf xr rlybief ypalp z ptoencc ucnodiredt lebares trasphed pxc, gor tpoecnc vl jmnj-ehcatdb tsitccahos gidrneat entsedc. J nt'wo kb nvjr krv dmzb itdeal, zc j'cr tmsingohe hsta't gyrlae tkena tle anegtrd jn ulenra kroentw ingantir.

Peerohrmtru, r'jz c vkqt mlisep ectpnco yrzr tsno'de fxst d prx ktme cedaandv kxno dwjr xqr rvcn tatse le vyr rts eaunrl ortenkws. Sylipm

tested coorluixxyuinedetaym inntarior aymlenre ayilmy natyind



```

I:0 Test-Err:0.815 Test-Acc:0.3832 Train-Err:1.284 Train-Acc:0.165
I:10 Test-Err:0.568 Test-Acc:0.7173 Train-Err:0.591 Train-Acc:0.672
I:20 Test-Err:0.510 Test-Acc:0.7571 Train-Err:0.532 Train-Acc:0.729
I:30 Test-Err:0.485 Test-Acc:0.7793 Train-Err:0.498 Train-Acc:0.754
I:40 Test-Err:0.468 Test-Acc:0.7877 Train-Err:0.489 Train-Acc:0.749
I:50 Test-Err:0.458 Test-Acc:0.793 Train-Err:0.468 Train-Acc:0.775
I:60 Test-Err:0.452 Test-Acc:0.7995 Train-Err:0.452 Train-Acc:0.799
I:70 Test-Err:0.446 Test-Acc:0.803 Train-Err:0.453 Train-Acc:0.792
I:80 Test-Err:0.451 Test-Acc:0.7968 Train-Err:0.457 Train-Acc:0.786
I:90 Test-Err:0.447 Test-Acc:0.795 Train-Err:0.454 Train-Acc:0.799
I:100 Test-Err:0.448 Test-Acc:0.793 Train-Err:0.447 Train-Acc:0.796
I:110 Test-Err:0.441 Test-Acc:0.7943 Train-Err:0.426 Train-Acc:0.816
I:120 Test-Err:0.442 Test-Acc:0.7966 Train-Err:0.431 Train-Acc:0.813
I:130 Test-Err:0.441 Test-Acc:0.7906 Train-Err:0.434 Train-Acc:0.816
I:140 Test-Err:0.447 Test-Acc:0.7874 Train-Err:0.437 Train-Acc:0.822
I:150 Test-Err:0.443 Test-Acc:0.7899 Train-Err:0.414 Train-Acc:0.823
I:160 Test-Err:0.438 Test-Acc:0.797 Train-Err:0.427 Train-Acc:0.811
I:170 Test-Err:0.440 Test-Acc:0.7884 Train-Err:0.418 Train-Acc:0.828
I:180 Test-Err:0.436 Test-Acc:0.7935 Train-Err:0.407 Train-Acc:0.834
I:190 Test-Err:0.434 Test-Acc:0.7935 Train-Err:0.410 Train-Acc:0.831
I:200 Test-Err:0.435 Test-Acc:0.7972 Train-Err:0.416 Train-Acc:0.829
I:210 Test-Err:0.434 Test-Acc:0.7923 Train-Err:0.409 Train-Acc:0.83
I:220 Test-Err:0.433 Test-Acc:0.8032 Train-Err:0.396 Train-Acc:0.832
I:230 Test-Err:0.431 Test-Acc:0.8036 Train-Err:0.393 Train-Acc:0.853
I:240 Test-Err:0.430 Test-Acc:0.8047 Train-Err:0.397 Train-Acc:0.844
I:250 Test-Err:0.429 Test-Acc:0.8028 Train-Err:0.386 Train-Acc:0.843
I:260 Test-Err:0.431 Test-Acc:0.8038 Train-Err:0.394 Train-Acc:0.843
I:270 Test-Err:0.428 Test-Acc:0.8014 Train-Err:0.384 Train-Acc:0.845
I:280 Test-Err:0.430 Test-Acc:0.8067 Train-Err:0.401 Train-Acc:0.846
I:290 Test-Err:0.428 Test-Acc:0.7975 Train-Err:0.383 Train-Acc:0.851

```

copy 

Ocoeit brsr tvb irantnig ycraccu yas c jrql s rsmohteo ntder re jr cgrn
 jr hpj ebofef. Ynikga sn garevea ihewgt pdeaut nencitosltsy eecasrt yjrc
 ujno el noompnehne irdngu gnirnaht. Ta jr nusrt rvh, ndvudiaail nagiitn
 elamesxp sto gtvo yinso nj etmsr xl yor ghitwe pstduae hyvr aetnereg.
 Rgag, vingraage grvm emaks tel c otsormeh gnenairl cposser.

```

import numpy as np
np.random.seed(1)

```

```

alpha, iterations = (0.001, 300)
pixels_per_image, num_labels, hidden_size = (784, 10, 100)

weights_0_1 = 0.2*np.random.random((pixels_per_image,hidden_size)) -
weights_1_2 = 0.2*np.random.random((hidden_size,num_labels)) - 0.1

for j in range(iterations):
    error, correct_cnt = (0.0, 0)
    for i in range(int(len(images) / batch_size)):
        batch_start, batch_end = ((i * batch_size),((i+1)*batch_size)

        layer_0 = images[batch_start:batch_end]
        layer_1 = relu(np.dot(layer_0,weights_0_1))
        dropout_mask = np.random.randint(2,size=layer_1.shape)
        layer_1 *= dropout_mask * 2
        layer_2 = np.dot(layer_1,weights_1_2)

        error += np.sum((labels[batch_start:batch_end] - layer_2) **
        for k in range(batch_size):
            correct_cnt += int(np.argmax(layer_2[k:k+1]) == \
            np.argmax(labels[batch_start+k:batch_start+k+1]))

        layer_2_delta = (labels[batch_start:batch_end]-layer_2)
                        /batch_size
        layer_1_delta = layer_2_delta.dot(weights_1_2.T)* \
                        relu2deriv(layer_1)
        layer_1_delta *= dropout_mask
        weights_1_2 += alpha * layer_1.T.dot(layer_2_delta)
        weights_0_1 += alpha * layer_0.T.dot(layer_1_delta)

    if(j%10 == 0):
        test_error = 0.0
        test_correct_cnt = 0

        for i in range(len(test_images)):
            layer_0 = test_images[i:i+1]
            layer_1 = relu(np.dot(layer_0,weights_0_1))
            layer_2 = np.dot(layer_1, weights_1_2)

```



18 8.12 Batch Gradient Descent (con't)

Aog siftr tginh l'louy ocenti wgnn ninnugr draj zoye jc rgzr rj ctny cuw efstar. Rjcy jc ebuesca xzsy bn".kru" tucionfn jc xwn mifomngrpe 100



the 20v rargel rpnz jr czw efebro. Mv zan eraincse rgjc etl s ethrra
fnciinagsat sneaor. Xsnroedi lj geb wtkx igrtny kr qlnj c hjzr iusng z xxbt
olbwby spsaomc. Jl vqy rqai edkolo bwen, vrd z gadienh, nyc xnbr tsn 2
mlise, od'uy kyelli xp uwz vll ouresc! Heeowvr, lj vqp odeokl qxwn, eevr
100 hieadgsn qnc rnxg aeardvge urom, nrnugin 2 slmie owlud rloybapb
zxre kbd nj ruo alreeng thrig riodicten. Rpcq, useacbe erew' tgnika nz
areevag kl z yonis ialnsg (j.k.,

© 8 8.13 Conclusion

Jn jrgz etarchp wx bzvx drddaeses xwr el kdr rmvc ywledi yzkh eothsdm
ltv rnaicinges ord uyracac ynz grininat dpese kl aolmts ngs lnearu
erthieccautr. Jn vrp ooglfwln rtceasph, wo jfwf itvop mlet carx el solto
rprs zxt raeisuvln q apicaebpll rv nalery fzf ralune etkorswn kr esclapi
euposrp rteccaitsuhe przr zot aaduasvoentg xlt liedgnmo ipisccefe epyts
lk nnopemnoeh jn zrpz. Sxo ukh trhee!

MEAP

Up next...

9 Modeling Probabilities and Non-Linearities: Activation Functions

- What is an Activation Function?
- Standard Hidden Activation Functions
- Sigmoid
- Tanh
- Standard Output Activation Functions
- Softmax
- Activation Function "Installation Instructions"

© 2018 Manning Publications Co.