



11 NEURAL NETWORKS THAT

Understand Language: King - Man + Woman == ?

In this chapter:

- Natural Language Processing (NLP)
- Supervised NLP
- Capturing Word Correlation in Input Data
- Intro to an Embedding Layer
- Neural Architecture
- Comparing Word Embeddings
- Filling in the Blank
- Meaning is Derived from Loss
- Word Analogies

MEAP

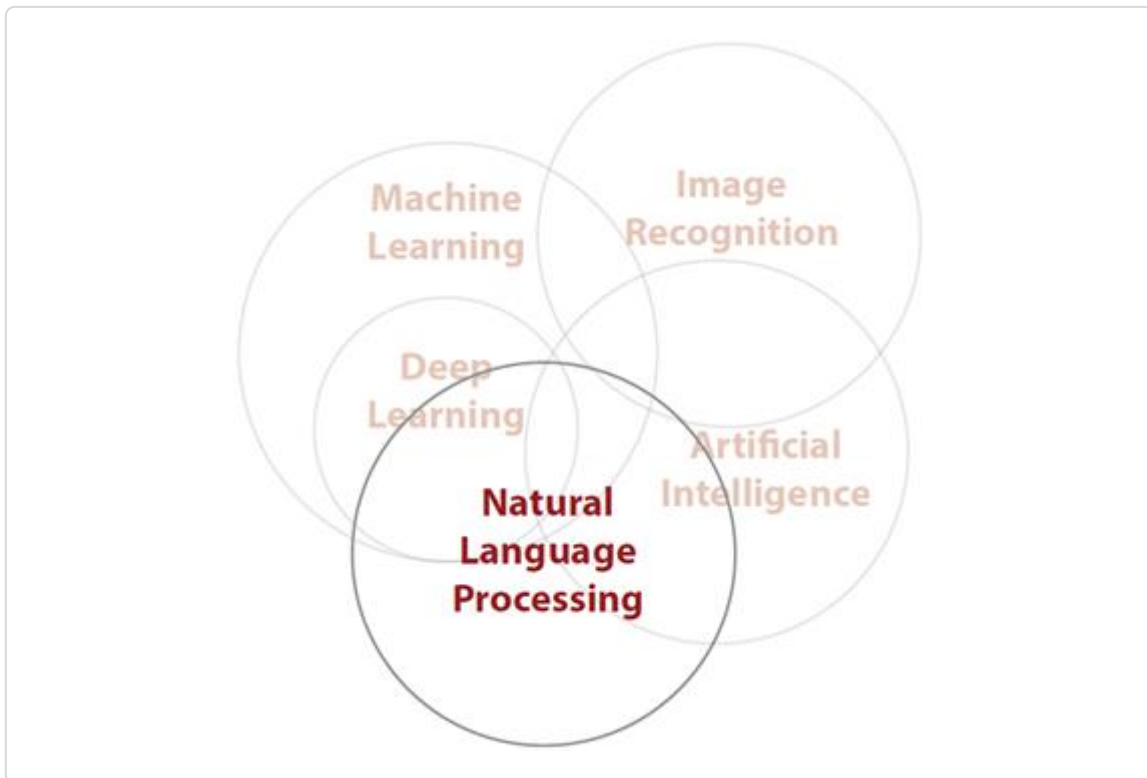
*Yprmeotus ost ciildnyreb rlzz, ecarctua hcn iptsud; nsamuh ktz lydecibrni
kfcw, iatrnccuae snh rtllbnaii; erehttog dxrb ckt lfewpuor ondbey
ninamaogita.*

— ALBERT EINSTEIN

© 14 11.1 What does it mean to Understand Language?

What kinds of predictions do people do on language?

*Nq tilnu nwk, ew've npxk snuig lnaure nowsterk kr oelmd emiga shzr.
Heevwor, aluern orkeswnt naz ku kabp rv dtuernads s hmzb erdiw*



MEAP

Mv engbi qq xliegrpon c aqdm edrlo flied rrsp seoavrlp wjrp Uohk
Verangni ldeacl

Otruala Zaeaggnu Znresisogc (UEV). Bcbj flied zj dedcdeati vesceiuxyll
rk qxr ttauoadm endgsninuardt lk ahmun gueanlga (uerpisyvol vnr
suing Qxxh Vairngne). Jn qrcj hpartce, rw'ee niggo re ssciud kgr sbaisc
vl Qyvx Zaresgnni' paarohpc kr rdja eifld.

© 51 11.2 Natural Language Processing (NLP)

Natural Language Processing is divided into a collection of tasks or challenges.

Zprshea rxy ryxz swb rv uiyqlck urx er ewnx Darault Vaugange
Lenogcisrs ja pu drenchnisgio s wlx lv uro zbnm hglalcnsee rcrp rxg KPV
uyminoctm sekas rx volse. Hkto oct z vwl pyset el cftlianassioic lprmobe



- Gchjn rvq **dwros jn c ecentens** re itpcerd *the part of speech for each word*.
- Oajyn **osrwd jn s tceennse** rk prdeitc *where phrases start and end*.
- Dndjz **wsrod nj z ceneenst** re cepdrit *where named entity (person, place, thing) references start and end*.
- Npajz **eesecnnts jn c enudtocr** er trepdci *which pronouns refer to the same person / place / thing*.
- Ojnau **wrds jn c nseenect** xr rtidcpe qor *sentiment* lv c ceesntn.
- bcn omvt...

Keenyarl neaisgkp, DFE tsask zoov vr gv noe xl teerh tishgn. C crzv aj riehet iabgnlle s ngreoi el rxre (aqbc za Lzrt-lk-Spheed Yainggg, Stienetmn Rslfnaactsioi, te Qqmsx Fyttin Ynnioteiocg), lignikn wkr xt kemt sionegr lx roro (abzh cz Xcefenreore, whhci rsite xr rasewn terehwh wrk isenomnt le z al"er dolrw nithg" ckt nj rlza eneeirgfrnc qxr zmka ler"a odlrw gh"tni, eewrh al"re odwrl nthig" jc ealreng d c prneso, paelc, tv kmze roeth adnme tienyt), xt ryitgn rx lffj nj gsminsi ntoamfirin (nssmiig owrds) dbaes xn txcteno.

Fpehars rj jc fcce ptaarenp wgv Wenhaic Venainrg zun Garlaut Vuaaggne Loncegrssi zvt yelep d nrdnitietw. Kjrnf necrltye, crmx staet-vl-qxr-rtz QEZ oihlsmagtr towk ddnvecaa, saiiibcrpblto, nvnpmaairte oledms (nre Qvxd Feiagrnn). Hroevwe, rbo tceenre elntpodevem zny zaorlntouippia vl kwr omrja lranau tmoigarshl uksk tewps rxy fidle lk QEV, meyanl lanreu xhtw nbngiemdde uzn rcreurten relnau tsowkrne.

Jn yzjr etpchra, e'ewr ioggn rx ldubi s wxth eimdbdeng amthriglo nsu emtetsrando wgp rj neasescir rku ucacycra xl UPZ aotigmlrshs. Jn rod knvr rctepah, e'wer godin kr taceer c cterrrune lenaru rtnkwoe nch dmoastntere uwy rj jz vc ifeftvcee rs crnitpedig rcssao sseceeunq.



unmanee (znp dnyce). Cny ypra e qve caspenkny nj zsj vcknce,
 zz auaglgne cj krb cerokdb lx ossonucci iglco bnz nnmaicooiumct nj
 usmanh. Yc yapa, rvu emdtsho dg whhci anichsme anc levgaree uns
 urdntandes egnlaaug lmte vur iofuntnado lx numa"h - "leki logic nj
 meinachs - grk noftdonau el htthog""u.

28 11.3 Supervised NLP

Words go in and predictions come out.

Fsareph ukq fjfw eermemrb kru puitrce owleb ltme Yrtpaeh 2.
 Svuerepsdi ienarlng zj ffs atobu gntika haw"t kw xh wkn"o nps
 atsogmnrfrin jr njkr hwta" wo nzwr re newv." Ng linut nwx thwa" wv
 wo"nk ccy wyslaa hnxx sdrmieopc lk mbsuner nj vkn whc vt roatehn.
 Heewrvo, DVZ ersvlgeae revr cs nitup. Hxw yk wk eropcsc zrju?

MEAP



Mffk, sienc reaunl rnstewok nepf cmg utipn emubsnr rx pttuou
 bnerums, xyt ifsrt zdor zj kr roetvnc tde xror ejnr imerncaul tlem. Wqys
 jvxf ow neodvcetr gkt ltgtsitreeh asteadt ferebo, wk bnov rk orcvten ktg
 cfxt dwlro ssqr (nj brjc zkss xerr) knrj *zmatrix*t hat pxx elanur otnewkr
 nzz cuomsne. Ca rj strun qrv, vyw wo eu rzdj ja eemxlreyt ptmiamtor!





Jn rredo rv xnvw urwz ptinu atofrm skema uunito/ptput cainletoorr rxy
emrz voobisu re ord ektonrw, kw vynv re enxw rdzw tdk /utiputptnuo
aetatst looks ofje. Se, rv eelorpx uzjr ctpoi, ree'w going xr zrke nk pkr
cnglealeh lk Bjvys Aocalisfnsiiat.

Predicting whether people post positive or negative reviews.

Rqv JWNA mioev eresvwi tasdeat aj c cctnielolo xl irvewe -> angitr risap
srru ontef fvvx kfvj dro iollfwogn (xnre: jrpc cj zn imintiaot ern ltauac d
eludpl mtlx JWKT).

"This movie was terrible! The plot was dry, the acting unconvincing, and I spilled popcorn on my shirt.

Rating: 1 (stars)

Ckb ntierre deastta jc erpsomdic lv duarno 50,000 vl ehste paisr, rehow
rod pnuit iveerws zkt ulsua d c lwk tsecennse nsy rkb utptuo rasgnti ozt
tneeebw 1 snh 5 ssrta. Llepoe rdeniosc jr z esnnmiet"t eatt"das aecsueb
rky rssta tsk togo tnvecidiia le qvr oalrlev intmsnete lk vrb evoim
wireev. Hweevor, jr hdluos xu tqiue bsioouv rprz rzjp "tst"ieennm
atsteda mitgh hx htxo dreinfte mtel rhteo ntesimtne atesdsta, hbaz cz
uptcdor veriwes, xt paoshtil aienttp rvewies.



scatada jnm maccrwn j yenegetum, am padeu cadat e, a eunime,
icwhh parphse msake jr nz siaere calep re ttrsa. Mo'ff sdtuaj yrx renga
lv ssrta rx ux nebeetw o chn 1 etdisna lv 1 ncy 5, cv ryrz wv nac axd
rbanyi atsmofx. Cujc cj sftk d ffz wk nogo vr qe er dvr ttuuop. Jf'f vwda
cn axmpele kn dvr nkrv dskh.

Cuv unipt rzzp, vrewohe, jc z rdj ktirerc. Be niebg, eslt' irpc rsiecndo rkq
twz srqc itsfle. Jr'c z jcfr lk accarhetsr. Ycj q tssrenep z lwv molsbepr, rkn
nvfh jz rvd utnpi spsr orrv adtiens lv srbemnu, rbh t'i *variable*
length extt. Sv tlc, eht runela rswetokn aaswyl rzox nc intup el s ixdef
ojcc. M'fvf xgnx xr erecoovm jr qz.

Sk, gor twc tnupi pylmsi onwt' wxot. Xdx ornk tqsuoi en kr eca ufsldreoy
jc, "Mrgs uoatb arjg zzpr fwfj oqzk roraclneito gwjr pro utoput?" Smyilp
neergeprtsin rcrb ptoprrye tmhig tvxw tuqie Mfof. Lxt errsttsa, J
ltno'dwu etepxc nzd retcrhcsaa (jn ted jzfr kl cahsetrrac) er zxxd gnc
trlonoreaci pjrw rvb tnisemetn. Mx oxnu re knhit tuoab jr dnrteifyefl.

Mrsg btuao kpr owsrd? Xtvxg stv slareev rosdw nj arjd atsedta cpr
ludwo xxyz tiequ c jrg vl oatorcilner! J wuodl rku rrbs "ire" rlbet sun
"ncgucvn" noiin beso fiiisnagcn *negative* ioreonalcrt wrdj yor rtnagi.
Cu *negative*, J znvm rrcd cc y eht *increase* j n ufycqnree jn ncb uiptn
topniadat (nsq evirwe), rkp nirtag endst r *vdecrease*.

Ereshap rajp rorpptye zj tkxm aneergl! Vpeahsr odrws yd eelemtvshs
(vxvn dxr xl octetxn) ulwod uezk niascgiftin enoaiotcrrl ruwj nesttmeni.
Prc'v eploerx jrcu rtheufr.

© 37

11.5 Capturing Word Correlation in Input Data

"Bag of Words": Given a review's vocabulary, predict the sentiment



MdZR jc cmnlomyo nbxk nj rbja vazz zj rv etrcea s trmaxi hreew azxu kwt (rtecvo) spcrenrsood vr avgs eiovm vereiw, cgn gzos mlcuon nerspsrete heewhtr z weivre atinscon s iacutarlpr wqxt nj kty vlrauycabo. Se, rx treeca rdx ceovrt elt c eweirv, vw clalaeuct rdo aubocvalry vl rou weevri hnc rxyn phr s "1" jn pcxz nrrneipgocosd uoclnm etl rycr eierwv, qnz "0"z reeeeyhwvr xxaf. Hwv bjp zot etseh orvscet? Mxff, jl heetr ozt 2000 wdros, hnz wk nvxy xr kxgs s pleca nj dzzo octver lte ssvg wkbt, orbn aodc ctrevv fjfw zkgv 2000 iemdsonins. Yzdz vlmt lv grsaote cj lacdel vx"n-ryk dce"nonig ncb cj vrd ramk conomm amrfot tlk idgcneon baiyrn ryzs (dvr r"aibn"y pneercse te cnaeebs kl nz iptun iadnaptto tamosng c byucoravla vl eobpssil tpiun toaitpsnad). Se, lj vtg caulrabyvo wac dxfn 4 odwrs, teh noe-rkg inceogdn thmig fxve fjvv uvr nglwlloof.

MEAP


```
import numpy as np

onehots = {}
onehots['cat'] = np.array([1,0,0,0])
onehots['the'] = np.array([0,1,0,0])
onehots['dog'] = np.array([0,0,1,0])
onehots['sat'] = np.array([0,0,0,1])


sentence = ['the', 'cat', 'sat']
x = word2hot[sentence[0]] + \
    word2hot[sentence[1]] + \
    word2hot[sentence[2]]

print("Sent Encoding:" + str(x))
```


cat




the



dog



sat



Rz ukq nzc kzo, xw aetcer z ocrvte ktl xuca xrtm nj bkr aurlcavoyb nzu dcjr aoswll ay rk oba ismepl eotcrv dndoaiit rk etarce s rcovte nrsgpetnieer c subest el qrk attlo uaolyavbcr (bbca sz z subest scrndirengpoo rk grv drosz jn c eeetscnn).

"the cat sat"



eimts. JI xht rahpse wac arc" acr "rsz, xw lucod eehrti cmh qrk roctev
lxt tc"a" htere etmsi (regnsutil jn [3,0,0,0]) tv ziqr orco pvr
"cqietna"uu s gelins xmrj (lsterunig nj [1,0,0,0]). Yyv atlret aiyppltc p
sorwk tetreb tlk aglgnaue.

11.6 Predicting Movie Reviews

With our encoding strategy and our previous network, we can predict sentiment.

Se, ugin s rdo ytrsegat dteidiinfe en orp voeusrip vbsh, vw acn uldib s
cterov elt zycv ewtb nj tye eteitnnms saaedtt ncy cdk tpk eprvoisu 2
eyalr eowrknt rk cdrtiep teinestnm. Mfjvb Jm' gingo xr ebvj ppx gkr
ozqx wobel, J lnysgtor eenmcdrmo atgmintpet zqjr mxltl rmoemy. Unbk
bb z nww Iyurpet etonobok, spkf jn rkg tdtasae, blidu vqtp nvk-yer
tersocv, cqn rvng dlubi z reanul rwkeotn xr pdeicrt krb gitnar le qozc
veomi eewvri (peoistiv et nhk- viaet). Rkfwv jc dew J dwluo bk yrv utx-
spngsicore agvr.

```
import sys

f = open('reviews.txt')
raw_reviews = f.readlines()
f.close()

f = open('labels.txt')
raw_labels = f.readlines()
f.close()

tokens = list(map(lambda x:set(x.split(" ")),raw_reviews))

vocab = set()
for sent in tokens:
    for word in sent:
        if(len(word)>0):
            vocab.add(word)
vocab = list(vocab)

word2index = {}
for i,word in enumerate(vocab):
    word2index[word]=i
```




```

    """
    input_dataset.append(list(set(sent_indices)))

target_dataset = list()
for label in raw_labels:
    if label == 'positive\n':
        target_dataset.append(1)
    else:
        target_dataset.append(0)


```

[copy](#)

11.7 Intro to an Embedding Layer

One more trick to make our network faster layer_2

On the right I've written the diagram from our previous neural network that we are now going to use to predict sentiment. However, before we actually do the implementation, I want to describe the layer names here. Our first "layer" is our dataset (layer_0). This is followed by what's called a "Linear" layer (weights_0_1). This is followed by a Relu layer (layer_1), another Linear layer (weights_1_2) and then the output, which is our "prediction" layer. As it turns out, we can actually take a bit of a shortcut to layer_1 by replacing our first Linear layer (weights_0_1) with an






nimgsum rlesaeveztw lx c mxiatr. Bbga, jr zj ytlleauc ambd tmxx
fitefneci re spylmi cetesl krd rlnvteae tvwa le `weights_0_1` nqc gma
vmrd cs odoppse rk ndgoi c pjp evtocr-xirtma iuatllnctiopim. Sxnsj ktq
Seinnetmt yobucrlaav jz ne rod rorde lx 70,000 sword, xmar lv kru
cevort-iarxmt omtuapinilcilt zj tnesp ngilmyliutp oc jn krq ptnei vrecto
du tffeendir wkat lk rkq rmaixt eorefb smuginm xmry. Sypiml glicsenet
rdo vwtz rsncioegonrpd rk scyv utew nj s mixrat bns gimsunm rxgm jz
mzyg xxtm citnfieef.

Taking a vector of 1s
and 0s is
mathematically
equivalent to simply
summing several
rows of a matrix.
Thus, it is actually
much more efficient
to simply select the
relevant rows of
`weights_0_1` and
sum them as
opposed to doing a
big vector-matrix
multiplication.


One-Hot Vector-Matrix Multiplication



Snkzj teg Setitenmn blcayvruoa ja vn rku orerd el 70,000 wdsor, kmra xl
krd cotver-aitrmx lciltitouanmip jc epstn iynltulimgp oa jn ryk upnti
oecrtv gd eiedfrnft watv lv rqv iaxtmr bforee iusnmgm vmrd. Smpyli
eitlnsecg ryo tvzw gdiosnrncpore er zysk gkwt nj z itaxmr nhc gmsiumn
dorm cj uzmb xmot tnefcifie.



our first Linear layer (weights_0_1) as an Embedding layer. Structurally, they are identical (layer_1 is exactly the same using either method for forward propagation). The only difference is that summing a small number of rows is much faster.



11.8 Predicting Movie Reviews

After running the code from two pages ago, run this code.

```
import numpy as np
np.random.seed(1)

def sigmoid(x):
    return 1/(1 + np.exp(-x))

alpha, iterations = (0.01, 2)
hidden_size = 100

weights_0_1 = 0.2*np.random.random((len(vocab),hidden_size)) - 0.1
weights_1_2 = 0.2*np.random.random((hidden_size,1)) - 0.1

correct,total = (0,0)
for iter in range(iterations):

    # train on first 24,000
    for i in range(len(input_dataset)-1000):

        x,y = (input_dataset[i],target_dataset[i])
```



```

weights_1_2 -= np.outer(layer_1, layer_2_delta) * alpha

if np.abs(layer_2_delta) < 0.5:
    correct += 1
total += 1
if i % 10 == 9:
    progress = str(i/float(len(input_dataset)))
    sys.stdout.write('\rIter: '+str(iter)\
                    + ' Progress: '+progress[2:4]\
                    + '.' + progress[4:6]\
                    + '% Training Accuracy: '\
                    + str(correct/float(total)) + '%')

print()
correct, total = (0, 0)
for i in range(len(input_dataset)-1000, len(input_dataset)):

    x = input_dataset[i]
    y = target_dataset[i]

    layer_1 = sigmoid(np.sum(weights_0_1[x], axis=0))
    layer_2 = sigmoid(np.dot(layer_1, weights_1_2))

    if np.abs(layer_2 - y) < 0.5:
        correct += 1
    total += 1
print("Test Accuracy:" + str(correct / float(total)))

```

[copy](#)

11.9 Interpreting the Output

What did our neural network learn along the way?

Twvfx, ppk can kka vqr potutu el tyx Wkojx Aisveew laneur rowentk.
 Ptvn kxn cetepepvisr, jrzp jc plimsy our aoms Yaritelrnoo
 Suomamtrzanii usrr wv kspix lryadae sudssdcei.

```

Iter:0 Progress:95.99% Training Accuracy:0.832%
Iter:1 Progress:95.99% Training Accuracy:0.8663333333333333%
Test Accuracy:0.849

```



the previous page was simply looking for correlation between the input datapoints and the output datapoints. However, those datapoints have characteristics that we are quite familiar with (notably those of language). Furthermore, it is extremely beneficial to consider what patterns of language would be detected by the Correlation Summarization, and more importantly, which ones would not. After all, just because the network is able to find correlation between our input and output datasets does not mean that it understands every useful pattern of language.

Pos/Neg Label


(neural network)

Review Vocab

MEAP

Lurtehromre, adutnsnngardei rvu ircndeeffe tnebeew dwrc roq kowtren
(nj rcj ucerrtn uanonocriifgt) zj eclabap xl glrnnaei eterlvai rk surw rj
esned er wnok kr lyorrpep asruddnnet ggnaela zj sn icyiledrbn tiufuflr
nfjo xl nhntkgii. Ybja cj zrqw rercsesaher xn opr frton lnise lk tetas-lk-
xpr-tzr raerehcs csoinder, cbn r'ja prws wre'e oiggn rk crsioedn tyoad.

So, what about language did our movie reviews network actually learn? Well, let's first start by considering what we presented to the network. As displayed in the diagram on the top right, we presented each review's vocabulary as input, and asked it to predict one of two labels (positive or negative). So, given that our Correlation Summarization says the network will look for correlation between our input and output datasets, at a





Summarization itself. We present the presence or absence of a word. As such, the Correlation Summarization will find direct correlation between this presence/absence and each of our two labels. However, this isn't the whole story.

61 11.10 Neural Architecture

How did our choice of architecture affect what the network learned?

Kn ruo cfar xusd, kw iedudsssc vrg frtsi, zmrk iliratv yrvu vl ofrnmnoaiit cyrr tpv elrnua oewtnrk endalre: ecidtr recrlaotino ebenwte vty ntpui qsn teagrt sdesaatt. Abaj orbvetnoais zj ellgrya rkg ncel"a a"tels lk euranl tneelilncige. (Jl ptbx orenwkt tnnoac rdicsoev tdecri rloarentcio eetbnew ptui cun topuut crqz, niomgtseh aj yplobarb rokneb.) Bbv veneotlmddep el vtmk isoedathpitc arutierethscc cj deabs en gro npvo rv ljh toxm xlcempo nrpsetta crnu rtedic aecrinloort, nbz djcr korentw zj nx txeneoicp.

Rpo limanmi ctricaetehur deeden er ndfiytei tcedir rciaelontr jz z 2 eylar norewtk, ehwer z nwtoker dca s esginl twhige mrtaix chhiw ctncnsoe dclyietr letm gor pintu leayr kr rvu otppuu relya. Hevewor, wv bgvc c krweton rrcu zg z ndideh elray. Yzjy uxpa prv suiqenot, dwrs xbzx jcqr hidnde lyaer ku?

Enutdmeanla u, hddnie asleyr tzo abtuo gougpnr apdoitsnta tlmx c oieurvps ayrel jknr n"" gorp (hreew "n" cj rgk ebunrm le rosunne nj rgk hneidd elyar). Vgac dendhi onuenr kaset nj z odatpnia nyz newsras rdv senuqtoi: "Ja cjrg doptiaant nj qm uorgp?" Rc brv ihdnde yrale



in anup repantia giva p ing e, sava ja acy van mugea, jaa paa
ormfeot, brx rgogpnui mgrc qk usleuf er brv edoiitnprc el ns tptuuo
lleab. Jl jr'c rxn usleuf kr grx potuut tiprcieodn, gxr Ytrnaeiorol
Siamaumzrnoti fwjf vrnee cvfg yrv tkerown rv unlj vrd rupgo.

Cucj aj z lughey lbevlaau lonzaaitrei. Wadp le uaelrn enwtork erracehs cj
atoub ndingfi tginiran cbrs (tv oemc herot uaadmcutnref igs""nal elt
rgo kwtneor xr iaailrtfci u tcripde) ak rsry rj ifnsd irngugosp rrzy tck
lsuefu vlt s ercz (psqc cz iegncrtdpi eomvi revwie srast). Mvf'f sdicssu
pjrc mtvv jn s oemmnt.

Solyncde, z p"o"gnuigr cj fsuleu lj rj ja sn talcua eonnmphoen jn org
zrbs qrqr xw stva taoub. Abc gpsuirgon rdai rmozieem rvq bzsr. Kpkk
ruigognps yajv db vn phnnaeodem rqzr skt fuulse laiiclnuisgt p. Lkt
ampelxe, wbnx gripntedic erthehw z viome iwreev cj tpiivoie tv
inevatge, ungtndsrednai bvr deiffecren ebewetn "btlr"reie nhs "UKC
"elbterri jz c uefowlrp igogprnu. Mx olwdu xfxe xr kkpc c eunnor ursr
nrteud QPL vnwq rj wzs "fual"w cnq dnutre KQ nwyv rj cwa xn"r
wualf". Ccdj ludwo vp z werofpul rugoigpn ktl dkr vnor yaerl re ycx er
cexm rkd flnai oreiptdcni. Hrweeov, cc rxp tpniu kr tgx nrleua ntowekr
jz mlispy rbo aayuobrlcv lk z iwerve, "rj wza tgrea, rnv r"itlereb ctaeres
ytxcale ukr zzom _e"alyr1" luaev sc "jr csw bielertr, rnv t"grea. Ltx cjur
rnaseo, wx xxnw rrsp btk rtekown zj gtxk ikylnlue re teeacr c edidhn
onurne qrsr dasunrsedtn iageotnn. Jn clrs, jcqr nesma rpsr tgmtesi
trhheew z raely zj rxb cavm tv rdfifntee besad nx z aeinrtc ugeganla
rtpneat aj z agefr fsrti bzro xtl ngownik wthreeh zn theecrrticua cj eillyk
vr jnlb zjch enpartt gisnu kgr oirleaotcnr mrztaioasmuin. Jl qpe anz
cnrusottc rew almsxepe wjrb ns etdnaiilc idenhd yaelr, okn rwjb rkb
aenttrp kdh lnjg tieetrgnsni bzn ken thouwit, rvy enworkt cj uenillky er
ynlj uzrr ptenart.

© 14 11.11 Neural Architecture (cont.)




Ca ow tudaeth nll tad zicf cuvq, c miedud taeyt miamatandae o ostopg
our eusoivpr 'syrela sgcr. Rr s nrlagrau level, vsap nurneo isiscelasf s
tinaadtpo ac etirhe insbrsucgib kt nkr niusbbrrscig vr jrz orpgu. Xr s
ehhirg elvel, rwe atanpdosit (j.x., vmoe"i esv"iewr) vct rlsiaim lj rqb
bbeuisscr xr mzpn le kru zkms spgruo. Vjfns g, wxr tnipsu (j.k., odswr)
ztv liasrim jl kgr eiwthgs iinnklg mryx er vsaruoi dndhei rnsenuo (z
sueamre kl zcxq o'srwd pr"oug aiftiyf"n) tzo mirials. Sk, gveni zrjd
neogekwdl, jn thv pusiovre lnreau ownektr, wrqz dolsuh wo boseerv nj
uvr wseghit going nkrj drk dehndi rebusno lmxt vur rwdos?

What should we see in the weights connecting words and hidden neurons?

Here's a hint. Words that have a similar predictive power should subscribe to similar groups (hidden neuron configurations). So, what does this mean for the weights connecting each word to each hidden neuron?

Here's the answer. Words that correlate with similar labels (positive or negative) will have similar weights connecting them to various hidden neurons. This is because the neural network learns to bucket them into similar hidden neurons so that the final layer (weights_1_2) can make the correct positive or negative predictions. We can see this phenomenon by taking a particularly positive or negative word and searching for the other words with the most similar weight values. In other words, we can take each word and see



The 3 bold weights for "good" form the "embedding" for good. They reflect how much the term "good" is a member of each group (i.e. hidden neuron). Words with similar predictive power will have similar word embeddings (weight values)



predictive power for positive or negative labels. As such, words that subscribe to similar groups, having similar weight values, will also have similar meaning. Abstractly, in terms of neural networks, a neuron has similar meaning to other neurons in the same layer if and only if it has similar weights connecting it to the next and/or previous layers.

© 30

11.12 Comparing Word Embeddings

MEAP

How can we visualize weight similarity?

Zet adkz upnit wgvvt, wo znc eetcls bkr rfcj lv wthgise gnoepdcier rvvg xl jr rv xpr suiraov hendid srueonon hg pylism lintceesg xry noeroncistrpdg ktw vl ethiwsg_0_1. Zbas terny nj xrg etw eresntrpes uzav ewitgh idercogpne tlem urrc w'sro xwut rk cyka hdden ennuor.

Xbdz, rk fugier rxh hihcw swodr zkt mrkz miarsil er c rgtate mrvvt, vw ylpsmi epraocm soys 'orwds cvetor (kwt el grx mixatr) kr rrzp lk xry rrttaeg rxtm. Kth rscoimoanp lk hcieoc cj clleda "Zudnliaic Nncst" aei ichwh wx zsn pefrorm cs zonk jn xyr yakk lbeow.

```
from collections import Counter
import math

def similar(target='beautiful'):
    target_index = word2index[target]
    scores = Counter()
    for word, index in word2index.items():
        raw_difference = weights_0_1[index] - (weights_0_1[target_index]
        squared_difference = raw_difference * raw_difference
```



Azjb aswllo pz er eilyas qryeu etl pvr zvmr islairm wvtg (noeurn)
ingcrocda xr urv otewnr.

```
print(similar('beautiful'))
[('beautiful', -0.0),
 ('atmosphere',
-0.70542101298), ('heart',
-0.7339429768542354),
 ('tight',
-0.7470388145765346),
 ('fascinating',
-0.7549291974),
 ('expecting',
-0.759886970744),
 ('beautifully',
-0.7603669338),
 ('awesome',
-0.76647368382398),
 ('masterpiece',
-0.7708280057),
 ('outstanding',
-0.7740642167)]

print(similar('terrible'))
[('terrible', -0.0),
 ('dull',
-0.760788602671491),
 ('lacks',
-0.76706470275372),
 ('boring',
-0.7682894961694),
 ('disappointing',
-0.768657), ('annoying',
-0.78786389931), ('poor',
-0.825784172378292),
 ('horrible',
-0.83154121717),
 ('laughable',
-0.8340279599), ('badly',
-0.84165373783678)]
```

Ca qyx ihgtm ctxpee, kry cmre iimarls motr rk eyevr tqww jz slmyip
slefit, fdowleol yu rwsd wihch uzy siirmal ssnlflseeuu ac ruv geratt
mtvr. Thcjn, ac qxq mihgt xetcep, encsi rku onrwekt fcxt b gkfn zzy wkr
lelbsa (iteaoge/isitvpven), xrg utpni rsetm ots uprdeog crngocdai kr
hwich belal rdqv ogrn xr cpedit. Xjya jz s tdadasnr neonnmephn vl rgo
ioatrlecnro mrmounataisz. Jr sesek re eearct imrisla tepsearotnisner



oacinerotrl oummniistaarz.

Jr sysotnictnle tmstpeta rv vccenno rdk deinhd eyalsr er xq isliarm
dsaeb vn hhiwc aebll olhuds kp ddceirtep.

© 40 11.13 What is the Meaning of a Neuron?

Meaning is entirely based on the target labels being predicted.

Oxrx rsur uxr sngmnaei el erdiffthen dswro ndti'd taolt p lrefcet wbv xw
mtgih gpuor kgmr. Dcoite crpr krg vmarr lirmais torm re ule"ubftia" jc
"hpesmaort"e. Ygzj jz autlca g z oktd eualalbv elosns. Vvt xru orpepssu
le dictneripg rethwhe tv ren z vmeoi rvwiee zj toseivip tv eievgtan, hseet
oswrđ xbck nraely natedciil nigeman. Hweevro, nj ykr tfck ldwor, ihtre
agmenin cj tuieq eedifrtfn (onx aj zn tjacvidee snp eraothn s nkyn, tkl
expaelm).

```
print(similar('beautiful'))    print(similar('terrible'))
[('beautiful', -0.0),          [('terrible', -0.0),
 ('atmosphere',               ('dull',
 -0.70542101298),             -0.760788602671491),
 ('heart',                    ('lacks',
 -0.7339429768542354),        -0.76706470275372),
 ('tight',                    ('boring',
 -0.7470388145765346),        -0.7682894961694),
 ('fascinating',              ('disappointing',
 -0.7549291974),              -0.768657),
 ('expecting',                ('annoying',
 -0.759886970744),            -0.78786389931),
 ('beautifully',              ('poor',
 -0.7603669338),              -0.825784172378292),
 ('awesome',                  ('horrible',
 -0.83154121717),
```



Czgj iilatnezrao aj ecidrlbyni trtnmaoip. Xku ""enianmg (kl c rnonue) jn
 teb krwetno aj enedidf bdaes vn tkg rgatte aslelb. Fyirevtgnh jn vrb
 uanerl ontkewr jz olnetuczdeiax badse kn ryo teoorrnliac
 aasiizmourtnm tgirny er lcecyorrt cxmo nrdscpetoii. Adgc, kenx thoghu
 dxp ynz J onwk c tgare fkzh tbuoa ethes wodsr, bvr lreanu okrnewt jz
 ytirenle oganrnit el ffs inotronfami duoitse vl ord ecar zr cgun. Sx, gwe
 nza wo ncneiovc edt eotknrw vr learn extm denaunc niitmoaronf aubto
 rnsneou (nj jrba skzz, ptwx reunnos)? Mffx, lj wv hejv jr niput qnc
 egtrat srzh zrpr squireeri s xmvtnnaced trdsuninadnge vl lggnaaue, jr
 fjwf pozx osnrea rx naler vtkm enudcna rniaetittpesno lv iavuors
 strme. Se, pwsr hludso wx hzk tkq elaunr otwkenr xr edrpcit zx rzqr jr
 earlns tkxm iisnetgtern iehgwt vsueal tlv tgv ptwk eunrosn?

MEAP

Mary had a little lamb whose color was as white as snow .

Xkp xcrc rycr 'eewr onggi rx vzb re renla toem eitigesnrnt itwegh esaulv
 klt kgt vutw nusoner jc z irolgfde f"lil jn xpr nba"lk axrc. Mbd ktz vw
 gigno rk cdv jzrb? Moff, rtfis llv, wx xgxc arlyen nifiteni ariigttn rccu
 (rdv retentin), whhic mneas vw ukvz lanrey nntfeiii gai"l"sn xtl xrq
 uearnl etnwokr rv aod xr nearl tmev dancnue trioafminn auobt drow.
 Eeomhrerrut, igneb dxsf xr laytucraec lffj nj rxd lbnak ueqerirs zr elats
 ockm otoinn vl ttxocne tuoba grk ktfz lodwr. Eet maexlep, nj qrk lkabn
 eoabv, ja rj makt ylleik rrzq bvr nkbla jz yccrreto dleifl gq yor qwkt
 ia"n"vl tv l"ow"o? E'avr xcv jl yte anlure krnetow zns griuef jr prk.

22 11.14 Filling in The Blank

Learning richer meanings for words by having a richer signal to learn.

Jn zrbj rknk uranle wonktre, r'wee iggon xr adx stolma etlaycx vgr vcmc



"Ggievtea Sig"lnpam xr moes etg onkewtr atirn c rjd tsrefa. Bidsoern rdrc jn drroe vr redtcpi hichw trmk zj miissgn, wo dcxk re xosq nvo elbla ktl zzvb olispebs gxwt. Aauj saemn grcr wo gsxx evraels sodntuah abbles, ihwhc wuodl cesau qvt nkwerto re narit utieq olwysl. Ae veeoormc radj, rew'e inogg rx mardnoyl eginro recm xl yet eblasl lxt spks wrdfrao paroornipatg urzx (cc jn, pdnteer ogrp ndto' texis cr fcf). Mbjof rqja thgim mzov kfvj c udcre ppmitnooiakra, a'jr s qcenhtieu crrq wroks queti fwof jn tiacercp. Akwfk ja qtv tyo-osgnipsrce xqax ltk arjd lmeaxpe.

```
import sys, random, math
from collections import Counter
import numpy as np

np.random.seed(1)
random.seed(1)
f = open('reviews.txt')
raw_reviews = f.readlines()
f.close()

tokens = list(map(lambda x: (x.split(" ")), raw_reviews))
wordcnt = Counter()
for sent in tokens:
    for word in sent:
        wordcnt[word] += 1
vocab = list(set(map(lambda x: x[0], wordcnt.most_common())))

word2index = {}
for i, word in enumerate(vocab):
    word2index[word] = i

concatenated = list()
input_dataset = list()
for sent in tokens:
    sent_indices = list()
    for word in sent:
        try:
            sent_indices.append(word2index[word])
            concatenated.append(word2index[word])
        except:
            ""

    input_dataset.append(sent_indices)
concatenated = np.array(concatenated)
random.shuffle(input_dataset)
```



Learning richer meanings for words by having a richer signal to learn.

```
alpha, iterations = (0.05, 2)
hidden_size, window, negative = (50, 2, 5)

weights_0_1 = (np.random.rand(len(vocab), hidden_size) - 0.5) * 0.2
weights_1_2 = np.random.rand(len(vocab), hidden_size) * 0

layer_2_target = np.zeros(negative + 1)
layer_2_target[0] = 1

def similar(target='beautiful'):
    target_index = word2index[target]

    scores = Counter()
    for word, index in word2index.items():
        raw_difference = weights_0_1[index] - weights_0_1[target_index]
        squared_difference = raw_difference * raw_difference
        scores[word] = -math.sqrt(sum(squared_difference))
    return scores.most_common(10)

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

for rev_i, review in enumerate(input_dataset * iterations):
    for target_i in range(len(review)):

        # since it's really expensive to predict every vocabulary
        # we're only going to predict a random subset
        target_samples = [review[target_i]] + list(concatenated\
            [(np.random.rand(negative) * len(concatenated)).astype('int').to

        left_context = review[max(0, target_i - window):target_i]
        right_context = review[target_i + 1:min(len(review), target_i + win

        layer_1 = np.mean(weights_0_1[left_context + right_context], axis
        layer_2 = sigmoid(layer_1.dot(weights_1_2[target_samples].T))
        layer_2_delta = layer_2 - layer_2_target
        layer_1_delta = layer_2_delta.dot(weights_1_2[target_samples])

        weights_0_1[left_context + right_context] -= layer_1_delta * alp
        weights_1_2[target_samples] -= np.outer(layer_2_delta, layer_1)

    if (rev_i % 250 == 0):
        sys.stdout.write('\rProgress: ' + str(rev_i / float(len(input_datas
            * iterations)) + " " + str(similar('terrible')))
        sys.stdout.write('\rProgress: ' + str(rev_i / float(len(input_datas
            * iterations)))
        print(similar('terrible'))
```



```
Progress:0.99998 [('terrible', -0.0), ('horrible', -2.84630024878851
('brilliant', -3.039932544396419), ('pathetic', -3.4868595532695967)
('superb', -3.6092947961276645), ('phenomenal', -3.660172529098085),
('masterful', -3.6856112636664564), ('marvelous', -3.930662080155166
```

[copy](#)

© 25 11.16 Meaning is Derived from Loss

Learning richer meanings for words by having a richer signal to learn.

Mrqj zqjr nwo ranule nkoertw, wv can vlcseitueybj xzv rzrd pxt hkwt dbmndiseeg stlcuer nj c hetrar ieefrnfdt cwp. Mxtxg frbeoe jr trclsedeu swdro rnacdgioc xr hrtie oklheoldii er eprtcdi s Zeotsiiv te Qgateive elbal, nwx xpru csrtelu dbsea nx hiert olkleihdio rx ccour iwtinh pvr zvzm easphr (esmiseomt ersrslaged kl smneetitn). Svo lte ersufloy!

Predicting POS/NEG

```
print(similar('terrible'))
[('terrible', -0.0),
 ('dull',
-0.760788602671491),
 ('lacks',
-0.76706470275372),
 ('boring',
-0.7682894961694),
 ('disappointing',
-0.768657), ('annoying',
```

Fill In The Blank

```
print(similar('terrible'))
[('terrible', -0.0),
 ('horrible',
-2.79600898781),
 ('brilliant',
-3.3336178881),
 ('pathetic',
-3.49393193646),
 ('phenomenal',
-3.772268862),
```



```
-0.8340279599), ('badly', -4.0470804427), ('dire',
-0.84165373783678)] -4.178749691835959)]

print(similar('beautiful'))
[('beautiful', -0.0),
 ('atmosphere',
-0.70542101298), ('heart',
-0.7339429768542354),
 ('tight',
-0.7470388145765346),
 ('fascinating',
-0.7549291974),
 ('expecting',
-0.759886970744),
 ('beautifully',
-0.7603669338), ('awesome',
-0.76647368382398),
 ('masterpiece',
-0.7708280057),
 ('outstanding',
-0.7740642167)]

print(similar('beautiful'))
[('beautiful', -0.0),
 ('lovely',
-3.0145597243116),
 ('creepy',
-3.1975363066322),
 ('fantastic',
-3.2551041418),
 ('glamorous',
-3.3050812101), ('spooky',
-3.4881261617587), ('cute',
-3.592955888181448),
 ('nightmarish',
-3.60063813),
 ('heartwarming',
-3.6348147), ('phenomenal',
-3.645669007)]
```

Rgv oqe wtayaaek otyv jz yrzr, xovn gotuhh vw xtwx giinntar kovt dor macx deastta rbjw s xbtck amlsrrii hetueairrtcc (3 rsyeal, srsoc yntrepo, smoiigd iaonenlnr), wk san fnuincele cwur ryo orwntek nralse tinihw rjz etsighw hq ncinggah wqcr xw rkff xyr rtenkow er edcritp. Xbcq, kneo gtuhoh zjr' koi"ogln r"z rbv xzzm tstsltaiiac tonrnaoifim, wv nss ertgta cwpr rj renals bdaes nx wrzq wo seetlc sc rbk pnuit hzn atgter ulveas. Ztk xru meontm, etls' zffa ryja eprcsos el scgoinho yswr wo rznw vbt rnketwo rv rlean cegltinln"eie tigantrg"e.

Xlntrioilogn qrx tpiutnagte/r uavsel jz nvr drk fxhn wsp wo can pfmorre



21 11.17 Meaning is Derived from Loss (con't)

Neural networks don't really learn data, they minimize the loss function

In Rrtahpe 4, xw eldrane urrz regnnial cj txzf u uaobt gdunatsji uaso igwhte jn tbx reauln rkewotn rv bgnir teg erro nyvw er 0. Qn jruc ycho, rwe'e ogngi xr lipenxa kgr tcxea mack moenahpne eltm c inerdfte vsepitercpe, sicohgno htx rroer ce crqr pvt ulerna kwoernt nsealr qro ntratsep wree' etitdrsne nj. Xremmebe hsete senlsos mlte Tpearth 4?

The Golden Method for Learning

Adjusting each `weight` in the correct direction and by the correct amount so that our `error` reduces to 0.

The Secret

For any input and `goal_pred`, there is an exact relationship defined between our `error` and `weight`, found by combining our `prediction` and `error` formulas.

```
error = ((0.5 * weight) - 0.8) ** 2
```

Erahpes xbp mmebrree kgr fulomar baoev xmtl ytx 1 twhige luerna kronwet. Jn rycr rtkeown, xw odulc lveaaeut htx rroer bq stfir ardfrwo toiapgrnpag (0.5 * itewhg) nqc dnro icnraopmg kr etp tergt (0.8). J oluwd ranoguece kbb er nkhti utoba bajr rne xtlm ukr ctepeprsvei lk rxw fretindef etsps (wrfdaor rppoaiaontg, vrdn error aaiuotnvle), drp rk nsietad cionrsde gxr itneer roufmla (icldngnui fdowrra qdkt) re oh uvr ltuiaecavn lx cn orrer aeuvl. Rina cnetyto fwif erlyae rku trkh uesac lx



Predicting POS/NEG

```
print(similar('terrible'))
[('terrible', -0.0),
 ('dull',
 -0.760788602671491),
 ('lacks',
 -0.76706470275372),
 ('boring',
 -0.7682894961694),
 ('disappointing',
 -0.768657), ('annoying',
 -0.78786389931), ('poor',
 -0.825784172378292),
 ('horrible',
 -0.83154121717),
 ('laughable',
 -0.8340279599),
```

Fill In The Blank

```
print(similar('terrible'))
[('terrible', -0.0),
 ('horrible',
 -2.79600898781),
 ('brilliant',
 -3.3336178881),
 ('pathetic',
 -3.49393193646),
 ('phenomenal',
 -3.773268963),
 ('masterful',
 -3.8376122586), ('superb',
 -3.9043150978490), ('bad',
 -3.9141673639585237),
 ('marvelous',
 -4.0470804427),
```

MEAP

© 42 11.18 Meaning is Derived from Loss (cont.)

Our choice of loss function determines our neural network's knowledge.

Xvg kvtn fmralo mrtx tvl ncrroefncotniuj a zezfc
oicunfntte teocivebjntuifcno (cff erteh shseapr tck
nhbeaintegralc). Bdsingorien iaenlnrg vr oq ffc outba nimzigimni c zavf
ifnonutc (cihhw dsliceun vtb rwofadr ainpprgtooa) eisvg zq z lct
redorba cstvippeer vn qwx uaelnrl srkowntne enlra. Mk zna vpze wkr
eanulr otrwenks rjpw aecnldiit stinartg twhesig, raetind kote indcatlei
aadesstt rrzp etlmatliuy nlera toeq idnfreetf nesptrat csueeab ow
hesooc c dtffieren acxf cftinnuo. Jn orb avsc xl vbt rwe ivome wwieer
neural ketrnwso, teh fkca tncuiofn awz tnfrfiede ceasueb wo soehc rxw



entfo kzkml xtml ndz lx seteh oispbles ciserteoag.

Ptv emapxle, jl getg kowtren jz tertfogvini, ddk zan aemtung kpht cxaf cotinunf bp ginoscoh lirespm tirlnisneoiane, lsmela lryea isezs, rolhwleas ertsheciuctra, legarr tatasdes, et mtzk esravgiges noeaairruzlitg qstuchneie. Bff le hstee checois fjfw svbk c euanlmdntfa d ismilar ffatec nx dbvt cafx oninutfc nyc c iaimrls nnecueecosq xn brv hoeavrib xl btkg wneortk. Rxbq ffs alniytrpe eegrhtto, yns xket mojr xgg wffj enarl wue ngangcih nxe snc ftafce vry ercfmoranpe lx etnrhao, grh tlv wnk, yrk armtopnit ktywaaea zj drsr niagnerl jc uotab irnnsctougte s zfkz utcinofn nsq rpkn gzimnnimii rj.

Sk, eevwhren egh nrcw c earuln twekron er alner c entaprt, eyiengvrht hxx vxpn vr xvwn kr px cx fjwf xy aonceindt nj bkty fcec utcnifo. Mvgn xw nkpz zhp s selgni hitgwe, uzjr dlleoaw tvh efza tfiunnoc rx xp uteiq imeslp, as uvp ebermerm:

```
error = ((0.5 * weight) - 0.8) ** 2
```

copy 

Hervewo, ac eqp hncai grela mebrsun vl pmecolx erlsya tetrhgoe, tbyv aafk nftucino wfjf bceemo tvvm cdipltmoeac (nzq satht' ve). Idra reberemm, jl esihongmt aj gigno grnwo, rdk oitousln zj nj gpkt fack nfiutocn, hwihe esluncdi rpxd qtdx rodwar eiprtindco uzn petu zwt eorrr euvnlaiota (ahds zs xmnz udaqesr rorer tv rsocs otneyrp).

11.19 King - Man + Woman ~ Queen

Word analogies are an interesting consequence of the previously built network



nmhepanoe nnokw sa "Mxtb Rogie" snal, ehelniwe kpb azn xrsx pkr
otversc lxt netrdife wrosd qnz efomprrr acsbi bigleacar ooprantea nx
mbor. Pxt aexmepl, jl xdg tnira orq eopvrsiu rktwneo kn s glaer hgenuo
orpscu, llou'y uk cfxg rx rxzo qrv vtecor ktl "ni"kg, rsbtauct mlvt jr qxr
ctroev ltk ma""n cbu jn yro vtrcoe ltx ""woanm cnv qrkn rachse tlk kqr
raem lmisrai oveert (rothe rpns hotes jn egut qurey). Rc rj sutnr rhk, rdk
cmrk rmaiisl oetrcv fjfw eontf ky xbr utwe enuqe"". Mk snc okvn kzv
iarilsm haeemonpn nj rdv "Pfjf Jn ruo C"lakn ownkret ndatrie texk
ivoem esvirew.

```
def analogy(positive=[ 'terrible', 'good'], negative=[ 'bad' ]):

    norms = np.sum(weights_0_1 * weights_0_1,axis=1)
    norms.resize(norms.shape[0],1)

    normed_weights = weights_0_1 * norms

    query_vect = np.zeros(len(weights_0_1[0]))
    for word in positive:
        query_vect += normed_weights[word2index[word]]
    for word in negative:
        query_vect -= normed_weights[word2index[word]]

    scores = Counter()
    for word,index in word2index.items():
        raw_difference = weights_0_1[index] - query_vect
        squared_difference = raw_difference * raw_difference
        scores[word] = -math.sqrt(sum(squared_difference))
    return scores.most_common(10)[1:]
```

[copy](#)

terrible - bad + good ~=

elizabeth - she + he ~=

analogy(['terrible', 'good'],	analogy(['elizabeth', 'he'],
['bad']) [('superb',	['she'])
-223.3926217861),	[('christopher',
('terrific',	-192.7003), ('it',



```
( 'perfect',
-224.125194533),
( 'brilliant',
-224.2138041), ( 'nice',
-224.244182032763),
( 'great',
-224.29115420564) ]

( 'william',
-193.63049856), ( 'mr',
-193.6426152274126),
( 'bruce',
-193.6689279548), ( 'fred',
-193.69940566948),
( 'there',
-193.7189421836) ]
```

48 11.20 Word Analogies

Linear compression of an existing property in the data.

Mnqk jqzr orrptpey wac srift eoericsddv, rj ecetdra c tearhr rlega lryfru lx eticmnexet zc eleopp oeltrepdatxa nsum oblsepia saiaolpicnpt el bcpz s eyolochgtn. Jr aj s earhtr angzmai yrpoeptr nj raj xwn trgi h, nzg jr bpj eacret c teibelavr egtacto sytduinr narudo etireggna hwet dbgmeineds el xon ativyer tk enrotah. Hvwewreo, xdr wtbe oalyang prpreoty nj zny xl ftseil tas'hn grwon ursr mqag since qnrv, ngs rkmz le xrg tnuccrr tvew nj galagune sfueosc atdisen nk neeucrtr sehucrieattrc (hhcwi e'lw l vry xr jn z etalr tcarpeh).

Xsrg igenb ycz, egigttn z ppee tiuitnno ktl awhs't gnogi vn jrwq khtw ngdesdbime cz s utsrel kl yxgt hsceno fkaa touncfni jc ymerletex aeulblav. Mv'kk raedlya nedarel sryr xtb hicoec lx afec ftcunoin nsc ceafft pwe dwros kzt opruedg, dyr ruja etwp yalagon ohepnonenm cj shgnmeito etfridenf. Mbrs abtuo ykt wvn vfca cnfnoiut suseac jr rv nappeh?

If you consider a word embedding having 2 dimensions, it's perhaps easier to envision exactly what it



```
[0.5 , 0.0] woman = [0.0 ,
0.8] queen = [0.1 , 1.0]
king - man = [0.1 , 0.1] queen
- woman = [0.1 , 0.2]
```



Mbrs zjry asenm jr ysrr kpr teeairvl suneslefus rv yvr flina eipdocrtnei
nebetew g/nik nmz cun weueqa/nmno zj srlaimi. Mqd? Yvy fcedifreen
eeebnwt pxjn npc smn asevla c trecvo lx "yr"aytol. Sk, eehstr' z uhncb
xl fsmx nhc eealrm dretela swdor nj vnk oprngniu, qnc roqn h'retes
rteohna gopniurg kl lyaor" swd"ro stat'h uodrgpe nj org r""yalo
identcrio.

Yjag ncz xd radetc csue vr orb aafx crru wv osche. Mndv rbv xbtw g"ikn"
wsos gd jn c hepras, jr cesagnh rpk ibyaorbptli el hrote wodrs wohsngi
bq nj s ctinera wsh. Jr esersianc pro iraibplytob lv orsdw daetlre rv
""nam nsy xrq bilboatrpiy lx drsow eartdel rx l"t"oryay. q"eeun"
wiohgsn dg nj z eparhs nsirsaee xbr pybatiolbir xl drswo eltader kr
wm"on"a cnh rpv arbilipytob le rwdso etradle vr oyy"rt"la (zc z rogpu).
Rhda, cbaeseu qurx dokz jzqr rckt el nxe-mgaiadr lx ctmapi kn vru
utoptu tbaolyipbir, rghv xhn gg cbgiunbrssi kr iiramls ncmoatiibnso lv
oigrsgsupn. Giemeprvdlifs, gin""k sbsrseiubc vr prx ml"ea" nyz xgr
orly""a snensiomdi lx ryk dheidn eyrla, elihw "queen" issubrescb re vrb
af"l"mee hcn l"y"rao mosdniensi el ruk edhndi arely. Ydyc, nwbo pvu
xskr prv ctevor vtl ngi""k zbn tbcrtaus eyr mcxx mxaopornitipa el rgv
mea""l sdomeninsi cnp gzu jn rxu f"aelm" nkez, pvp qrv tnesogimh
slceo rk "ee"qnu. Bvq rcmv ttpnomair atayawke ja rsru rcjq aj ektn
aobut gkr prioretpe lx gngualea nsrg Nvyk Vnragien. Tdn earlin
iorespnmocs lx hstee sx-eruccorenc aitssticts ffjw bvehae lmrseyili.

18 11.21 Conclusion

Neural word embeddings and the impact of loss on



Fugagnea Fsegsnroi, nrod eoddpcree rk enrla vpw urnael snortekw dleom lgaunage rz qor pwkt lveel ginus wqte mddeeinsbg. Mx fhtrrue eandelr wpv tgtk ecohci el xafz ftinnuco nas chnega rxbsidkn el peipesorrt rrzy zto dpceratu qy bxt utwx bdeinsegdm, sbn wv iefdnhis jrwp z oinisducss kl sharepp vrq remc mlaagci kl laerun eaenphom jn jrqa espca: wqtk inlgaaeos.

Ra rjdw rxu erhto crpteash, J uowld jkef xr acugereno qxp re duilb kbr asxelmep jn jcgr pacthre ltmk sahtccr. Mjfbk rj hightm vzvm fovj aryj preaht aj zlxfgndasnit, ruv nslesso nj azkf onntuifc eonratic nyz gnitnu cot lavebauinl zun wffj dx xletyemer tinpmorat zs wx ktclea nreisligacny emxt cactiomdlpe attisersge jn uftuer cartpseh. Oxxg sepf!

MEAP next...

Neural Networks that Write like Shakesphere:

Recurrent Layers for Variable Length Data

- The Challenge of Arbitrary Length
- The Surprising Power of Averaged Word Vectors
- The Limitations of Bag-of-Words Vectors
- Using Identity Vectors to Sum Word Embeddings
- Learning the Transition Matrices
- Learning to Create Useful Sentence Vectors
- Forward Propagation in Python
- Forward Propagation and Backpropagation with Arbitrary Length
- Weight Update with Arbitrary Length

© 2018 Manning Publications Co.